
Catching a Moving Subspace: Low-Rank Bandits Beyond Stationarity

Anonymous Author(s)

Affiliation

Address

email

Abstract

Low-rank linear bandits replace the ambient dimension d with an intrinsic rank $r \ll d$ in the regret rate, but only while the reward subspace stays fixed; once it moves, the rank advantage disappears and existing low-rank methods break. We study what it takes to *catch a moving subspace* from scalar rewards. Concretely, we consider linear contextual bandits over T rounds where the reward parameter factorizes as $\theta_t = B_k^* w_t$ for a piecewise-constant orthonormal rank- r factor $B_k^* \in \mathbb{R}^{d \times r}$ whose range $\text{range}(B_k^*)$ changes across an unknown number K of segments, with latent state $w_t \in \mathbb{R}^r$; the learner observes a scalar reward $y_t = x_t^\top \theta_t + \varepsilon_t$ for action $x_t \in \mathcal{A}_t \subset \mathbb{R}^d$ with noise ε_t , and may purchase information by designating rounds as *probes* at per-probe cost $c > 0$. Stationary low-rank methods cannot track changing subspaces; nonstationary methods ignore low-rank structure and pay the ambient rate $\tilde{O}(d\sqrt{T})$. We first establish an *identification boundary*: three probe-side conditions (known noise variance, bounded state-noise coupling, full-dimensional probe support) are individually necessary and jointly sufficient for recovering $\text{range}(B_k^*)$ from quadratic functionals of scalar rewards. We then introduce **SPSC** (Single-Play Subspace-Calibrated Optimism), which interleaves isotropic probes with ridge-UCB exploitation *inside* the learned r -dimensional subspace via a quadratic-measurement identity. The costed dynamic regret is $\tilde{O}(r\sqrt{T}) + \tilde{O}(T^{2/3}) + O(WV_{\text{in}})$, with W the exploitation window and V_{in} the within-segment path variation, replacing the ambient $d\sqrt{T}$ rate with intrinsic rank r . A CUSUM-style adaptive variant removes the oracle-boundary assumption at the same leading rate. On eleven benchmarks spanning synthetic, UCI/MovieLens, clinical, and production-log data, SPSC beats state-of-the-art nonstationary and low-rank baselines whenever $d - r \gtrsim T^{1/6}$, matching the analytical crossover.

1 Introduction

The promise of low-rank bandits is to replace the ambient dimension d with an intrinsic rank $r \ll d$ in the regret rate, but only while the underlying subspace stays fixed. Once it moves, neither stationary low-rank nor ambient nonstationary methods are equipped to track it. This paper asks how to *catch a moving subspace* efficiently, given only scalar rewards.

High-dimensional sequential decision problems, such as personalized recommendation, clinical dosing, and targeted scientific experimentation, share two properties that the linear-bandit literature has so far handled only in isolation. First, the reward parameter $\theta_t \in \mathbb{R}^d$ at round $t \in \{1, \dots, T\}$ (where T is the horizon and d is the ambient dimension) typically

admits a *low-rank* factorization $\theta_t = B_k^* w_t$ with intrinsic rank $r \ll d$, where $B_k^* \in \mathbb{R}^{d \times r}$ has orthonormal columns, $w_t \in \mathbb{R}^r$ is a latent state, and the segment index $k \in \{1, \dots, K\}$ ranges over $K \geq 1$ piecewise pieces of the horizon: user preferences and clinical response patterns typically live on a thin slice of the ambient space. Second, that thin slice is *not static*: tastes drift, cohorts change, treatment regimes get revised, and the subspace $\text{range}(B_k^*)$ itself moves across unknown change points $1 = \tau_0 < \tau_1 < \dots < \tau_K = T+1$ that delimit the segments $\mathcal{I}_k = [\tau_{k-1}, \tau_k)$. Crucially, the learner sees only a scalar reward $y_t = x_t^\top \theta_t + \varepsilon_t$ for an action $x_t \in \mathcal{A}_t \subset \mathbb{R}^d$ with noise $\varepsilon_t \in \mathbb{R}$, so the low-rank structure can neither be read off nor refreshed from direct observations of θ_t .

Both ingredients of this problem appear in real bandit deployments. Spotify’s homepage contextual bandit, deployed at scale, explicitly targets daily shifts in user intent across user cohorts [Feijer et al., 2025], production evidence that bandit decisions in deployment face moving structure rather than stationary parameters. The Open Bandit Dataset [Saito et al., 2021], a public benchmark we evaluate on, ships scalar bandit feedback collected across multiple ad campaigns at ZOZOTOWN. In our clinical benchmarks, the IWPC warfarin cohort itself exhibits cohort-segmented predictive structure: a recent re-analysis identifies three latent patient subgroups whose dose-response coefficients flip sign across boundaries [Liu et al., 2025], demonstrating in the very dataset we benchmark on that the rank- r effective subspace is not constant across the population. What is absent is a combined treatment in linear contextual bandits: rewards living on a low-rank subspace whose identity shifts at unknown change points.

Existing work treats exactly one of the two properties. Stationary low-rank bandits [Jun et al., 2019, Jedra et al., 2024, Lu et al., 2021] deliver the $\tilde{O}(r\sqrt{T})$ rate but assume a fixed B^* , which any nontrivial subspace change invalidates ($\tilde{O}$ suppresses polylog factors in $T, d, r, 1/\delta$). Nonstationary bandits [Russac et al., 2019, Cheung et al., 2019] adapt to drifting parameters but pay the full ambient rate $\tilde{O}(d\sqrt{T})$ and do not exploit intrinsic rank. Closing this gap is delicate: with scalar feedback, a single response y_t is a rank-one measurement of $\theta_t \theta_t^\top$, so the subspace can only be recovered through *second-order* probing, and every probe diverts a round away from exploitation, which must itself be priced. No prior algorithm recovers a changing subspace under scalar feedback, and no prior analysis prices that recovery against the exploitation regret it enables.

Within each segment, the latent state $w_t \in \mathbb{R}^r$ follows a stable linear dynamical system $w_t = A_k w_{t-1} + \eta_{t-1}$ ($A_k \in \mathbb{R}^{r \times r}$, innovation $\eta_{t-1} \in \mathbb{R}^r$). The learner may designate a subset of rounds $\mathcal{T}_{\text{probe}}$ as *probes*, playing $x_t = u_t \sim Q$ for a fixed probe distribution Q on $\mathbb{R}^d$ instead of an action in $\mathcal{A}_t$. Performance is the *costed dynamic regret*

$$\text{DynReg}_T^{(c)} := \sum_{t=1}^T \left(\max_{x \in \mathcal{A}_t} x^\top \theta_t - x_t^\top \theta_t \right) + c |\mathcal{T}_{\text{probe}}|, \quad (1)$$

penalizing both suboptimal actions and probe rounds (each at cost $c > 0$). With m_k probes per segment, subspace error shrinks as $O(1/\sqrt{m_k})$ while probe cost grows as $O(cm_k)$; the optimum of this tradeoff sets the leading rate.

Identifiability boundary. Three probe-side conditions, namely known noise variance, bounded state-noise coupling, and full-dimensional probe support, are each individually necessary (in the unrestricted-second-moment problem) and jointly sufficient for recovering $\text{range}(B_k^*)$ from single-play scalar rewards, via a lifted operator $\mathcal{K}^{-1}$ (the inverse of the probe moment operator $\mathcal{K}$, defined below) acting on the centered statistic $y_t^2 - \hat{\sigma}^2$ (with $\hat{\sigma}^2$ a noise-variance estimate) (Theorem 2.2; Propositions B.12–B.14). Dropping any one makes the predictable second moment $\widehat{M}_t := \mathbb{E}[\theta_t \theta_t^\top \mid \mathcal{H}_{t-1}]$ unidentifiable from quadratic measurements alone, irrespective of algorithm.

Contributions.

- 1. Identification boundary (Theorem 2.2).** We route probe-round feedback through a *lifted moment* $G_t := \mathcal{K}^{-1}(s_t u_t u_t^\top)$ with $s_t = y_t^2 - \hat{\sigma}^2$, where $\mathcal{K}(M) = \mathbb{E}[(u^\top M u) u u^\top]$

- 86 is the probe moment operator and $\hat{\sigma}^2$ is the noise-variance estimate. Theorem 2.2 es-
 87 tablishes *tight identifiability*: (a) under the three probe-side conditions, the segment-
 88 averaged $\hat{M}_k$ recovers $\text{range}(B_k^*)$ at rate $1/\sqrt{m_k}$; (b) each of the three conditions is
 89 individually necessary in the unrestricted second-moment problem, so the probe-side
 90 assumption set is minimal at the level of quadratic functionals of scalar rewards. The
 91 necessity half is what distinguishes this from a generic identifiability result: it identifies
 92 the boundary of the solvable region, not merely a sufficient interior point. Within the
 93 rank- r LDS class, rank structure may relax (i); (ii) and (iii) remain needed.
- 94 2. **Algorithm. SPSC** (Alg. 1) interleaves probing with windowed projected ridge-UCB
 95 inside the learned r -dimensional subspace; a **SPSC-Adaptive** variant (Alg. 2) uses a
 96 CUSUM-style detector to discover segment changes online with only a delay overhead
 97 relative to the oracle-boundary guarantee.
- 98 3. **Regret (Theorem 4.1)**. For the interleaved Alg. 1 with a balanced probe schedule,
 99 fixed K , and suppressing polylog factors, $\text{DynReg}_T^{(c)} \leq \tilde{O}(r\sqrt{T}) + \tilde{O}(T^{2/3}) + O(WV_{\text{in}})$,
 100 replacing the ambient $d\sqrt{T}$ rate with the intrinsic rank.
- 101 4. **Experiments**. We evaluate SPSC on eleven benchmarks (a synthetic phase-transition
 102 grid, five UCI environments plus MovieLens, two semi-synthetic clinical datasets, the
 103 Russac small- d benchmark, and Open Bandit production logs) against eleven baselines
 104 (LinUCB, D-LinUCB, SW-LinUCB, Restart-LinUCB, LowOFUL, VOFUL, LowRank-
 105 Reward, LinTS, SW-LinTS, and two adapted stationary low-rank methods BOSS and
 106 Jedra). SPSC delivers regret reductions whenever $d - r \gtrsim T^{1/6}$, qualitatively consis-
 107 tent with the bound. A direct *necessity* stress test (App. G.3, part C) restricts probe
 108 coverage to a proper subspace and observes the $\Omega(T)$ regret blow-up predicted by Propo-
 109 sition B.14, confirming that the third identifiability condition is not a proof artifact.

110 **Related work**. Stationary low-rank bandits [Jun et al., 2019, Kang et al., 2022, Jedra
 111 et al., 2024, Jang et al., 2024, Stojanovic et al., 2023, Duong et al., 2024, Lu et al., 2021]
 112 achieve intrinsic-rank rates but assume a fixed B^* ; nonstationary bandits [Russac et al.,
 113 2019, Cheung et al., 2019, Abbasi-Yadkori et al., 2023, Hou et al., 2024, Kim et al., 2022]
 114 bound dynamic regret in the ambient dimension and ignore low-rank structure; cost-aware
 115 observation [Seldin et al., 2014, Tucker et al., 2023, Elumar et al., 2024] prices information
 116 acquisition for fixed parameters, not for identifying a changing subspace. SPSC is the first
 117 algorithm we are aware of that (i) attains regret scaling with intrinsic rank r , (ii) adapts to
 118 changing subspaces without oracle boundaries, and (iii) prices probe acquisition explicitly.
 119 Appendix A gives the full positioning (Tables 5–7).

120 2 Setting and probe-based identification

121 **Piecewise low-rank model**. At every round t the reward parameter admits a factoriza-
 122 tion $\theta_t = B_k^* w_t$, where $B_k^* \in \mathbb{R}^{d \times r}$ has orthonormal columns and is piecewise constant over
 123 $K \geq 1$ segments $\mathcal{I}_k$ with unknown change points $1 = \tau_0 < \tau_1 < \dots < \tau_K = T+1$, and the
 124 latent state $w_t \in \mathbb{R}^r$ follows a stable linear dynamical system inside each segment

$$w_t = A_k w_{t-1} + \eta_{t-1}, \quad t \in \mathcal{I}_k,$$

125 with unknown A_k satisfying $\rho(A_k) \leq 1 - \alpha_0$ for some $\alpha_0 > 0$ (used only to ensure $\|w_t\| \leq S_w$
 126 via the stationary covariance bound) and zero-mean innovation η_{t-1} with covariance $\Sigma_{\eta,k} \succ$
 127 0. We assume uniform action-boundedness $\sup_{x \in \mathcal{A}_t} \|x\| \leq R_{\mathcal{A}}$, a known probe distribution
 128 Q with $\mathbb{E}[uu^\top] \succeq \rho_Q I_d$ and sub-Gaussian marginals (bounded $\|u\| \leq L$ is a special case; see
 129 Remark 2.1), conditionally σ_ε -sub-Gaussian noise with unknown variance, bounded state-
 130 noise coupling $\|\mathbb{E}[\varepsilon_t \theta_t]\|_2 \leq \epsilon_\times$, and full-dimensional probe coverage. Theorem 2.2 below
 131 shows these three probe-side conditions are individually necessary and jointly sufficient for
 132 identifying the changing subspace from single-play scalar rewards. Throughout, $\delta \in (0, 1)$
 133 denotes the global confidence parameter (concentration bounds hold with probability $\geq 1 - \delta$)
 134 and $\lambda > 0$ denotes the ridge regularizer used by the windowed estimator of §3.

135 **Remark 2.1** (Probe distribution: scaled sphere). *The theory uses scaled-sphere probes*
 136 $u_t = \sqrt{d} v_t$, $v_t \sim \text{Unif}(\mathbb{S}^{d-1})$, which satisfy $\|u_t\|_2 = \sqrt{d}$ exactly (so no truncation event

is needed) and $\mathbb{E}[uu^\top] = I_d$. In our implementation we draw $\tilde{u} \sim \mathcal{N}(0, I_d)$ and rescale $u \leftarrow \sqrt{d} \tilde{u} / \|\tilde{u}\|$, which is exactly the scaled-sphere distribution. The same argument applies to isotropic Gaussian probes $u_t \sim \mathcal{N}(0, I_d)$ via a χ_d^2 truncation at $L = \sqrt{2d \log(2T/\delta)}$ (Appendix C).

Quadratic measurement identity. Let $\hat{\sigma}^2$ be the algorithm's estimate of the noise variance and define the centered probe statistic $s_t := y_t^2 - \hat{\sigma}^2$ on every probe round $t \in \mathcal{T}_{\text{probe}}$. Under the probe-noise assumptions of §2, for $\delta_\sigma := \hat{\sigma}^2 - \sigma_\varepsilon^2$,

$$\mathbb{E}[s_t \mid \mathcal{H}_{t-1}, u_t] = u_t^\top \widetilde{M}_t u_t + 2u_t^\top \mathbb{E}[\varepsilon_t \theta_t \mid \mathcal{H}_{t-1}, u_t] - \delta_\sigma, \quad (2)$$

where $\widetilde{M}_t := \mathbb{E}[\theta_t \theta_t^\top \mid \mathcal{H}_{t-1}]$ is the predictable second moment, whose range is contained in $\text{range}(B_k^*)$ for $t \in \mathcal{I}_k$; under the non-degenerate innovation condition $\Sigma_{\eta,k} \succ 0$ (§2), the probe-time average has range exactly $\text{range}(B_k^*)$ (Proposition B.9). Under exact probe conditions ($\delta_\sigma = \epsilon_\times = 0$) this collapses to $\mathbb{E}[s_t \mid \cdot] = u_t^\top \widetilde{M}_t u_t$: a scalar reward carries a quadratic measurement of $\widetilde{M}_t$.

Lifted subspace estimator. The probe moment operator $\mathcal{K} : M \mapsto \mathbb{E}[(u^\top M u) u u^\top]$ admits a closed-form inverse on symmetric matrices for scaled-sphere probes $u = \sqrt{d} v$, $v \sim \text{Unif}(\mathbb{S}^{d-1})$:

$$\mathcal{K}(M) = \frac{d}{d+2} (\text{tr}(M) I_d + 2M), \quad \mathcal{K}^{-1}(N) = \frac{d+2}{2d} N - \frac{\text{tr}(N)}{2d} I_d, \quad \|\mathcal{K}^{-1}\|_{\text{op} \rightarrow \text{op}} \leq 1 \quad (3)$$

on the symmetric subspace (Lemma B.4). Define the *lifted probe sample*

$$G_t := \mathcal{K}^{-1}(s_t u_t u_t^\top), \quad \mathbb{E}[G_t \mid \mathcal{H}_{t-1}] = \widetilde{M}_t + \widetilde{B}_t, \quad (4)$$

with bias $\widetilde{B}_t = -(\delta_\sigma/d) I_d$ exactly on scaled-sphere probes, hence $\|\widetilde{B}_t\|_{\text{op}} = |\delta_\sigma|/d$ (Lemma B.6). Averaging lifted samples across the probe rounds of segment k yields $\widehat{M}_k := m_k^{-1} \sum_{t \in \mathcal{T}_k} G_t$ (with $\mathcal{T}_k := \mathcal{T}_{\text{probe}} \cap \mathcal{I}_k$ the probe rounds in segment k , $m_k := |\mathcal{T}_k|$), whose top- r eigenvectors form $\widehat{U}_k \in \mathbb{R}^{d \times r}$ with $\|\widehat{P}_k - P_k^*\|_{\text{op}} \leq C_{\text{sub}} \sqrt{\log(d/\delta)/m_k}$ (Corollary B.10 below). The variance misspecification bias $\widetilde{B}$ is a scaled identity (population level), shifting all eigenvalues uniformly and *not rotating eigenvectors*: subspace recovery is robust to the plug-in estimate of σ_ε^2 at the population level. Finite-sample projector concentration still has a floor $\Delta_\sigma \propto |\hat{\sigma}^2 - \sigma_\varepsilon^2|/(d \lambda_{\min})$ due to the random fluctuation around the biased mean (Cor. B.10).

Positive identification through the lifted moment G_t above and the necessity result that follows together yield the following theorem, the foundational result of the paper.

Theorem 2.2 (Identification under priced probing). *Consider the model of §2 with scaled-sphere probes $u_t = \sqrt{d} v_t$, $v_t \sim \text{Unif}(\mathbb{S}^{d-1})$, and the lifted probe sample $G_t = \mathcal{K}^{-1}(s_t u_t u_t^\top)$ from (4).*

1. **Identifiability.** Under (i) known noise variance σ_ε^2 ($\delta_\sigma = 0$), (ii) bounded state-noise coupling $\|\mathbb{E}[\varepsilon_t \theta_t]\|_2 \leq \epsilon_\times$, and (iii) full-dimensional probe support, the lifted sample satisfies $\mathbb{E}[G_t \mid \mathcal{H}_{t-1}] = \widetilde{M}_t$ exactly (the bias $\widetilde{B}_t = -(\delta_\sigma/d) I_d$ from Lemma B.6 vanishes under (i); the state-noise coupling term vanishes by sphere symmetry, regardless of $\epsilon_\times$), and the segment-averaged estimator $\widehat{M}_k = m_k^{-1} \sum_{t \in \mathcal{T}_k} G_t$ satisfies, with probability at least $1 - \delta$, the projector concentration $\|\widehat{P}_k - P_k^*\|_{\text{op}} \leq C_{\text{sub}} \sqrt{\log(d/\delta)/m_k}$; in particular $\widehat{U}_k = \text{TopEig}_r(\widehat{M}_k)$ recovers $\text{range}(B_k^*)$ at rate $1/\sqrt{m_k}$ (Proof in Appendix B.2).
2. **Necessity (unrestricted-second-moment scope).** Each of (i), (ii), (iii) is individually necessary in the unrestricted second-moment problem: removing any one yields two distinct conditional second moments $\widetilde{M}_t \neq \widetilde{M}'_t$ that generate identical observation laws $\mathbb{E}[s_t \mid \mathcal{H}_{t-1}, u]$, so $\widetilde{M}_t$ is not identifiable from quadratic measurements. The three constructions are independent (each counterexample satisfies the remaining two conditions), so the three form a minimal identifiability set in this class. Within the rank- r LDS class of §2 ($r < d$), rank deficiency partially relaxes (i) since σ_ε^2 is identifiable from the $d - r$ smallest eigenvalues; (ii) and (iii) remain needed (Proof in Appendix B.3).

3 Algorithm

SPSC (Alg. 1) alternates two phases:

- **Probe rounds** ($t \in \mathcal{T}_{\text{probe}}$): draw $u_t \sim Q$ (the probe distribution from §2; e.g. the scaled-sphere distribution of Remark 2.1), observe y_t , accumulate G_t in the current-segment estimator $\widehat{M}_t$, and refresh $\widehat{U}_t = \text{TopEig}_r(\widehat{M}_t)$.
- **Exploitation rounds** ($t \notin \mathcal{T}_{\text{probe}}$): project each action $x \in \mathcal{A}_t$ to $z_t(x) = \widehat{U}_{t-1}^\top x \in \mathbb{R}^r$, maintain a *windowed* ridge-UCB estimate $\widehat{a}_t = \widetilde{V}_t^{-1} \widetilde{b}_t$ over a sliding window $\mathcal{W}_t$ of the last W exploitation rounds, and play $x_t = \arg \max_{x \in \mathcal{A}_t} \{z_t(x)^\top \widehat{a}_t + \beta_t^{(r,W)} \|z_t(x)\|_{\widetilde{V}_t^{-1}} + \gamma_t \|x\|_2\}$.

The confidence radius $\beta_t^{(r,W)} = \sigma_\varepsilon \sqrt{r \log(1 + WR_{\mathcal{A}}^2/(\lambda r)) + 2 \log(2K/\delta)} + \sqrt{\lambda} S_w$ scales with $\sqrt{r}$, not $\sqrt{d}$: this is where the intrinsic-rank improvement enters. The correction term $\gamma_t = \Gamma_k + V_{k,t}(W)$ absorbs subspace-mismatch error $\Gamma_k \propto \varepsilon_k := \|\widehat{P}_k - P_k^\star\|_{\text{op}}$ and local within-window drift $V_{k,t}(W) := \sum_{s, s+1 \in \mathcal{I}_k, t-W \leq s < t} \|\theta_{s+1} - \theta_s\|_2$ (formal definitions in App. B.4). Per-round cost is $O(dr|\mathcal{A}_t| + r^2)$ vs. $O(d^2|\mathcal{A}_t|)$ for ambient LinUCB, an $O(d/r)$ speed-up in addition to the statistical gain.

Full pseudocode for Alg. 1 (oracle boundaries) is in Appendix D, alongside the adaptive variant described next.

Adaptive variant (unknown boundaries). SPSC-Adaptive removes the oracle-boundary assumption by detecting segment changes online: it maintains two non-overlapping rolling-window estimators $\widehat{M}_t^{\text{recent}}, \widehat{M}_t^{\text{past}}$ of the probe-time second moment and triggers a subspace reset whenever $S_t := \|\widehat{M}_t^{\text{recent}} - \widehat{M}_t^{\text{past}}\|_{\text{op}}$ exceeds a calibrated threshold b . The threshold is chosen so that, with probability $1 - \delta_{\text{FA}}$, no false alarm occurs across the horizon, while every true change of size $\Delta_k \geq 2b$ is detected within a bounded delay D_{max} (specified in Appendix D together with the full pseudocode, tuning, and CUSUM analysis).

4 Regret guarantees

Balanced probe schedule. The probe schedule is *balanced*: there is an absolute constant a_0 such that, for every segment $k \in \{1, \dots, K\}$ and every prefix-count $j \in \{0, 1, \dots, m_k\}$, the number of exploit rounds in E_k with exactly j probes collected before them is at most $a_0 \lceil \ell_k / m_k \rceil$. (Uniformly spaced probes satisfy this.)

Theorem 4.1 (Costed dynamic regret of Alg. 1). *Run Alg. 1 with oracle segment boundaries $\{\mathcal{I}_k\}_{k=1}^K$, scaled-sphere probes, and a balanced probe schedule. Under the model of §2, Assumption B.2, and exact centering $\widehat{\sigma}^2 = \sigma_\varepsilon^2$, set the per-segment probe budget to $m_k = \min\{\ell_k, \lceil c_0 \ell_k^{2/3} \rceil\}$ for an absolute constant $c_0 > 0$ (chosen as the closed-form optimum (14); segments shorter than the burn-in threshold $q^\star = O(\log(KdT/\delta)/\lambda_{\min}^2)$ are charged trivially), where $\lambda_{\min} > 0$ is the predictable probe-time r th-eigenvalue lower bound (Lem. B.1). Then with probability at least $1 - \delta$ and for fixed K ,*

$$\text{DynReg}_T^{(c)} \leq \widetilde{O}(r\sqrt{T}) + \widetilde{O}(T^{2/3}) + O(WV_{\text{in}}). \quad (5)$$

Including all constants: $\text{DynReg}_T^{(c)} \leq \widetilde{O}(r\sqrt{KT}) + \widetilde{O}(A^{1/3}B^{2/3}K^{1/3}T^{2/3}) + O(R_{\mathcal{A}}WV_{\text{in}}) + \widetilde{O}(K^{2/3}T^{1/3}/\lambda_{\min}^2)$, with $A := (R_{\mathcal{A}} + R_Q)S_w + c$, $B := CR_{\mathcal{A}}S_w(1 + R_{\mathcal{A}}\sqrt{W/\lambda})\sqrt{\log(2KdT/\delta)}$, $R_Q := \sup_{u \in \text{supp}(Q)} \|u\|_2$ ($R_Q = \sqrt{d}$ for scaled-sphere; $R_Q \leq R_{\mathcal{A}}$ if probes lie inside the action set); the $\widetilde{O}(K^{2/3}T^{1/3}/\lambda_{\min}^2)$ summand is the burn-in cost (lower order than $\widetilde{O}(T^{2/3})$ for fixed K). Here $V_{\text{in}} := \sum_k \sum_{s, s+1 \in \mathcal{I}_k} \|\theta_{s+1} - \theta_s\|_2$ is the within-segment path variation and W is the exploitation window (cumulative on E_k when $W \geq \ell_k$); the bound is most informative when V_{in} is small. Proof in §B.5.

Corollary 4.2 (Plug-in variance). *Estimate $\widehat{\sigma}^2$ from N probe-round residuals (Lem. B.11) and write $\delta_\sigma := |\widehat{\sigma}^2 - \sigma_\varepsilon^2|$. For scaled-sphere probes, the bias $\widetilde{B} = -(\delta_\sigma/d)I_d$ from Lem. B.6 is*

a scaled identity, hence preserves the eigenvectors of $\bar{M}_k^{\text{probe}}$. Davis–Kahan applied with the biased mean as reference removes the Δ_σ floor at population level, so Theorem 4.1 holds with the displayed bound, $\lambda_{\min}$ replaced by $\lambda_{\min} - \delta_\sigma/d$ in the constants. Choosing $N \gtrsim T^{2/3}/d^2$ yields $\delta_\sigma/d \ll \lambda_{\min}$ with high probability, so the leading rate is preserved.

Corollary 4.3 (Within-segment-stationary special case). *Suppose $V_{\text{in}} = 0$, i.e. within each segment $\mathcal{I}_k$ the parameter is constant ($\theta_t \equiv \theta^{(k)}$ for $t \in \mathcal{I}_k$, with $\theta^{(k)} \in \text{range}(B_k^*)$). The non-degenerate innovation assumption ($\Sigma_{\eta,k} \succ 0$) still holds across the segment family, so the rank- r subspace $\text{range}(B_k^*)$ is well-defined and identifiable through SPSC’s probe channel. The bound of Theorem 4.1 reduces to*

$$\text{DynReg}_T^{(c)} \leq \tilde{O}(r\sqrt{T}) + \tilde{O}(T^{2/3})$$

for fixed K (suppressing polylog factors). The leading $\tilde{O}(r\sqrt{T})$ term comes from LinUCB on the r -dimensional projected target $z_t^\top a^{(k)}$ (this does not require the within-segment target to itself span an r -dimensional set; the projector is what collapses the geometry to dimension r , not the trajectory of θ_t). This is the regime that class- or cohort-induced segment shifts approximate empirically.

Scope of the bound. Theorem 4.1 analyzes exactly the practical Alg. 1: probes are interleaved according to a balanced schedule, $\hat{U}_t$ is refreshed online, and exploitation rounds re-project the windowed history through the current basis. The time-varying-basis subspace error $\varepsilon_{k,t} = \|\hat{P}_t - P_k^*\|_{\text{op}}$ is controlled uniformly in t via a prefix concentration argument (Lemma B.17).

The leading $\tilde{O}(r\sqrt{T})$ term beats the ambient baseline $\tilde{O}(d\sqrt{T})$ whenever $r \ll d$; the $\tilde{O}(T^{2/3})$ term is the price of acquiring the subspace information online. The constant B depends on the exploitation window W as $\sqrt{W}$; for W treated as a problem parameter (the setting throughout this paper, consistent with $W \leq \min_k \ell_k$ on ℓ_k -bounded segments), this is absorbed into the constants and the $\tilde{O}(T^{2/3})$ rate is preserved. If W is allowed to scale with T , the second term becomes $\tilde{O}(W^{1/3}T^{2/3})$. Solving $r\sqrt{T} \lesssim T^{2/3}$ gives a simple rule of thumb: SPSC tends to outperform ambient LinUCB whenever $d - r \gtrsim T^{1/6}$. The experiments in §5 show a crossover qualitatively consistent with this rule of thumb.

Corollary 4.4 (Rank-adaptive guarantee). *If the true rank r is unknown, thresholding $\widehat{M}_k$ at $\tau_k^{\text{rank}} := 2R_X \sqrt{\log(2d/\delta)/m_k}$ (R_X the lifted-sample envelope from App. B.2) recovers the true rank with probability $\geq 1 - \delta$ whenever the population eigengap exceeds $4\tau_k^{\text{rank}}$, and the bound (5) applies unchanged. Proof in §B.6.*

Proposition 4.5 (Adaptive SPSC). *Suppose every true change satisfies $\Delta_k \geq 2b$ for the threshold b set in Appendix B.7. Then with probability at least $1 - \delta_{\text{FA}}$, SPSC-Adaptive (Alg. 2) attains the same leading rate as Theorem 4.1 with an additive $O(KW_{\text{det}})$ delay overhead. Proof in Appendix B.7.*

5 Experiments

We evaluate SPSC and SPSC-Adaptive on eleven benchmarks: a synthetic phase-transition grid; UCI/MovieLens (Covertime, Pendigits, Satimage, MNIST, Fashion-MNIST, MovieLens); clinical (Warfarin, Vancomycin); the Russac small- d stress test (§5.4); and Open Bandit production logs (§5.5). Baselines: LinUCB [Abbasi-Yadkori et al., 2011], D-LinUCB [Russac et al., 2019], SW-LinUCB [Cheung et al., 2019], Restart-LinUCB, LowRank-Reward, LowOFUL [Jun et al., 2019], VOFUL [Kim et al., 2022], LinTS, SW-LinTS, and adapted BOSS [Duong et al., 2024]/Jedra [Jedra et al., 2024]; Oracle-LinUCB (true subspace) is an optimistic reference. Curves are means over ≥ 10 seeds with ± 1 SE bands.

5.1 Phase-transition boundary

Theory predicts a crossover at $d - r \asymp T^{1/6}$: SPSC wins above it, ambient LinUCB below. We verify on a piecewise low-rank LDS at $T=5,000$, sweeping $(d, r) \in$

275 $\{5, 10, 20, 30, 45, 60, 80, 100\} \times \{1, 3, 5, 10, 15, 20\}$ with $r < d$ (40 cells), $K=10$, $\sigma_\varepsilon=0.3$, spectral
 276 radius 0.99, 40 actions, probe cost $c=0.1$.




Figure 1: **Phase-transition boundary.** SPSC/LinUCB costed-regret ratio across the (d, r) grid. The crossover $d - r = T^{1/6}$ (dashed) separates SPSC-wins (ratio < 1) from ambient-wins; per-rank curves track $\sqrt{r/d}$ (dotted), matching Theorem 4.1.

277 SPSC achieves 13–34% reductions whenever $d \geq 45$ with every r -curve tracking $\sqrt{r/d}$: use
 278 SPSC when $d - r \gtrsim T^{1/6}$, and ambient LinUCB otherwise.

279 5.2 Real-data multi-baseline comparison

280 Six real-data benchmarks recast as piecewise low-rank bandits on a (d, r) -grid: UCI **Cover-**
 281 **type** ($K=4$, $T=10,000$); **Pendigits**, **Satimage**, pixelized **MNIST** and **Fashion-MNIST**
 282 (random-projection features), and **MovieLens-100K** with genuine user ratings (all $K=10$,
 283 $T=5,000$, 10 seeds). Segments are class or user-subset shifts. The first five have cluster-
 284 induced low-rank structure; MovieLens is only approximately low-rank (rank- d SVD fea-
 285 tures, §E) and serves as a robustness stress test. Table 1 reports SPSC vs. LinUCB and vs.
 286 best non-oracle competitor on representative cells; per-cell tables in Appendix F.

Table 1: **SPSC/LinUCB** and **SPSC-Adaptive/Best Competitor** ratios across six real-data benchmarks (10 seeds per cell). Values <1 indicate SPSC wins; **bold** marks cells where an SPSC variant beats every non-oracle baseline. SPSC trails LinUCB at $d \approx r$ (top of each block) and gains as d grows relative to r . Restart-LinUCB, the closest competitor, is subspace-oblivious, so its regret is constant along each row.

Dataset / r	$d=55$		$d=105$		$d=155/200$		best	
	S/Lin	S/Best	S/Lin	S/Best	S/Lin	S/Best	cell	ratio
Covertypes, $r=10$	0.96	0.99	0.88	0.92	—	—	(105, 20)	0.86
Pendigits, $r=10$	0.76	0.77	0.65	0.73	—	—	(105, 10)	0.65
Satimage, $r=10$	0.47	0.59	0.40	0.58	—	—	(105, 10)	0.40
MNIST, $r=10$	0.84	0.86	0.74	0.88	0.73/0.98	0.98	(105, 20)	0.74
Fashion-MNIST, $r=10$	0.88	0.90	0.73	0.88	0.67/0.99	0.99	(105, 20)	0.71
MovieLens, $r=10$	1.00	1.00	0.97	0.97	0.92/0.93	0.93	(200, 5)	0.92

287 At $(d=105, r=10)$, SPSC-Adaptive beats LinUCB by -47% on Pendigits (operating-regime
 288 grid in Fig. 2, App. F) and -68% on Satimage (closest competitor VOFUL is $+112\%$ above
 289 SPSC-Adaptive; per-cell results in Table 10, App. F); at $(d=105, r=20)$ on Covertypes it
 290 beats Restart-LinUCB by -15% . Non-low-rank baselines are by construction r -insensitive.

291 5.3 Clinical benchmarks: Warfarin and Vancomycin dosing

292 **Warfarin.** Warfarin dosing depends on demographics, genotype (VKORC1, CYP2C9),
 293 and clinical indicators [Consortium, 2009]. We calibrate to the IWPC cohort ($\sim 5,700$ pa-

Table 2: **Warfarin** (representative cell $r=3$, mean $\pm$ SE, 10 seeds). SPSC-Adaptive cuts regret by 67.3% vs. LinUCB. Full per-rank tables in Appendix F.

Method	Costed regret	vs. LinUCB
Oracle LinUCB	177 ± 4	-91.2%
SPSC Alg. 1 (ours)	1372 ± 33	-31.6%
SPSC-Adaptive (ours)	657 ± 39	-67.3%
LowOFUL	683 ± 97	-66.0%
VOFUL	779 ± 93	-61.2%
LowRank-Reward	1056 ± 40	-47.4%
SW-LinUCB	2096 ± 18	+4.5%
LinUCB	2006 ± 17	—

tients, $d=93$, $K=8$ segments, $T=5,000$, 10 seeds; reward = negative deviation from therapeutic dose). Table 2 ($r=3$): SPSC-Adaptive cuts regret by 67.3% vs. LinUCB; full per-rank tables in Table 14 (App. F); random-subspace ablation in App. I.

Vancomycin. With the AUC-targeted Rybak et al. [2020] protocol (Cockcroft–Gault $\text{CrCl} \rightarrow \text{CL}_{\text{vanco}} = 0.04 \text{CrCl}$, daily dose at $\text{AUC}_{24}=500$; $d=93$, $K=8$, $T=5,000$, 10 seeds), Table 3 sweeps $r \in \{1, 2, 3, 5, 10\}$: SPSC-Adaptive cuts regret by 36.4% at $r=1$ and 14.5–25.7% for $r \in \{2, 3, 5, 10\}$ vs. LinUCB, beating every non-oracle baseline at every rank. Per-rank tables in Appendix K.

Table 3: **Vancomycin** (mean $\pm$ SE, 10 seeds): SPSC-Adaptive control regret across range of r vs. LinUCB and the best non-oracle baseline. Full tables in Appendix K.

r	SPSC-Adaptive	vs. LinUCB	vs. best comp.
1	525.3 ± 45.8	-36.4%	-35.1%
2	613.6 ± 33.1	-25.7%	-17.1%
3	640.3 ± 33.0	-22.4%	-10.1%
5	665.9 ± 27.2	-19.3%	-12.7%
10	705.6 ± 19.5	-14.5%	-7.0%

5.4 Standard nonstationary benchmark

On the small- d stress test of Russac et al. [2019] ($d=2$, $r=1$, $K=4$, $T=6,000$, 50 arms, 30 seeds), SPSC reduces regret by **-46.5%** vs. D-LinUCB, **-34.8%** vs. SW-LinUCB, and **-52.0%** vs. stationary OFUL (itself +11.5% above D-LinUCB here). Full table in Appendix H.

5.5 Real exploration logs: Open Bandit Pipeline (ZOZO)

The Open Bandit Dataset [Saito et al., 2021] ships production exploration logs (random-policy ZOZOTOWN slice, $T=5,000$, $K=10$, 10 seeds); only approximately low-rank, so this is a *robustness* stress test rather than a setting where the rank- r assumption holds exactly. At $d=55$, $r=5$, SPSC reaches **83 ± 2** vs. LinUCB’s 102 ± 1 (-18.8%); SPSC-Adaptive -10.9%. SPSC neither breaks nor over-claims on weakly-low-rank production data. Per-cell results in Table 15 (App. F).

5.6 Comparison with stationary low-rank methods (oracle restarts)

For a best-effort comparison with stationary low-rank methods, we adapt BOSS [Duong et al., 2024] and Jedra [Jedra et al., 2024] by restarting at every true segment boundary, an oracle advantage we do not give SPSC (setup in Appendix J). Table 4: SPSC wins all eight cells by 5.8–19.9%, with the lead largest at small d (-16.3 to -19.9% at $d=55$) and narrowing to 5.8–8.0% at $d=200$. SPSC-Adaptive (no oracle) also beats both on every cell.

Table 4: **Adapted stationary low-rank methods on a (d, r) -grid.** Costed regret (mean $\pm$ SE, 10 seeds, $T=5,000$, $K=10$). BOSS/Jedra restart at *oracle* segment boundaries; SPSC-Adaptive detects them online. SPSC variants win every cell.

d	r	SPSC Alg. 1	SPSC-Adaptive	BOSS (oracle)	Jedra (oracle)	SPSC vs. best comp.
55	5	323 $\pm$ 10	342 $\pm$ 15	399 $\pm$ 17	386 $\pm$ 18	−16.3%
55	10	471 $\pm$ 10	507 $\pm$ 13	588 $\pm$ 27	602 $\pm$ 24	−19.9%
105	5	264 $\pm$ 9	274 $\pm$ 10	300 $\pm$ 14	299 $\pm$ 13	−11.7%
105	10	377 $\pm$ 15	401 $\pm$ 14	423 $\pm$ 18	429 $\pm$ 14	−10.9%
105	20	578 $\pm$ 12	608 $\pm$ 14	654 $\pm$ 17	654 $\pm$ 14	−11.6%
200	5	208 $\pm$ 5	218 $\pm$ 6	236 $\pm$ 9	226 $\pm$ 9	− 8.0%
200	10	277 $\pm$ 10	292 $\pm$ 12	294 $\pm$ 11	305 $\pm$ 15	− 5.8%
200	20	442 $\pm$ 13	452 $\pm$ 11	473 $\pm$ 16	489 $\pm$ 18	− 6.6%

Robustness of SPSC-Adaptive. Detailed comparison under correct vs. misspecified oracle boundaries is in Appendix L: SPSC-Adaptive is robust to false-alarm boundaries, while Alg. 1 degrades monotonically as the oracle is corrupted.

5.7 Rank misspecification, probe-rate sensitivity, and take-away

Rank misspecification (Covertime, $d=155$, $r^*=10$; App. G.2): sweeping over nine grid points $r \in \{1, 3, 5, 10, 15, 20, 30, 50, 80\}$ gives a U-shaped curve (under: +33% at $r=1$; sweet spot $r=30$, −10%); SPSC beats LinUCB across $r \in [15, 50]$, so mild overestimation is safe. **Probe rate** (App. G.1): on the synthetic reference setting ($d=4$, $r=1$), the probe-period sweep $\{5, 10, 20, 30, 50, 100, 300\}$ has its optimum near 20–30, matching $m_k^* \propto \ell_k^{2/3}$. **Additional datasets** (see Appendix F): SPSC wins every cell at $d \geq 105$, peaking at $(d, r)=(105, 20)$. **Sensitivity & dynamics.** Large-scale assumption sweeps (App. G.4) confirm SPSC tracks Oracle within a constant factor through variance misspecification, approximate-rank perturbation ($\epsilon_k^\perp$ up to 0.4), and spectral radius up to $\rho=0.95$; SPSC remains the best non-Oracle method at $\rho=0.99$. Subspace error tracks the predicted $1/\sqrt{m}$ rate (App. G.5); noise/change-point-frequency/drift-speed sweeps (App. G.6) confirm the same trends. **Take-away.** Across the eleven benchmarks, SPSC delivers $\sqrt{r/d}$ intrinsic-rank savings whenever $d - r \gtrsim T^{1/6}$ and does not catastrophically fail outside that regime.

6 Conclusion

Sequential decision problems in recommendation, clinical dosing, and adaptive experimentation share two features: rewards live on a low-rank subspace of rank $r \ll d$, and that subspace shifts at unknown change points. Stationary low-rank, nonstationary, and cost-aware bandits each cover one piece; none recover a changing subspace under priced single-play scalar feedback. We close this gap: three probe-side conditions, known noise variance, bounded state-noise coupling, and full-dimensional probe support, are each individually necessary (in the unrestricted-second-moment scope) and jointly sufficient for subspace recovery from quadratic functionals of scalar rewards (Theorem 2.2). Based on this, we propose SPSC, which interleaves probing with windowed projected ridge-UCB in the learned r -dimensional subspace and achieves costed dynamic regret $\tilde{\mathcal{O}}(r\sqrt{T}) + \tilde{\mathcal{O}}(T^{2/3}) + O(WV_{\text{in}})$, replacing the ambient $d\sqrt{T}$ rate with the intrinsic r (Theorem 4.1); a CUSUM adaptive variant attains the same rate without oracle change points (Proposition 4.5). On eleven benchmarks across synthetic, UCI, clinical, and production data, SPSC reduces regret whenever $d - r \gtrsim T^{1/6}$, a crossover consistent with theory.

References

Yasin Abbasi-Yadkori, Dávid Pál, and Csaba Szepesvári. Improved algorithms for linear stochastic bandits. *Advances in neural information processing systems*, 24, 2011.

- 355 Yasin Abbasi-Yadkori, András György, and Nevena Lazić. A new look at dynamic regret for
356 non-stationary stochastic bandits. *Journal of Machine Learning Research*, 24(288):1–37,
357 2023.
- 358 Rajendra Bhatia. *Matrix analysis*. Springer Science & Business Media, 2013.
- 359 Wang Chi Cheung, David Simchi-Levi, and Ruihao Zhu. Learning to optimize under non-
360 stationarity. In *The 22nd International Conference on Artificial Intelligence and Statistics*,
361 pages 1079–1087. PMLR, 2019.
- 362 International Warfarin Pharmacogenetics Consortium. Estimation of the warfarin dose with
363 clinical and pharmacogenetic data. *New England Journal of Medicine*, 360(8):753–764,
364 2009.
- 365 Thang Duong, Zhi Wang, and Chicheng Zhang. Beyond task diversity: provable represen-
366 tation transfer for sequential multitask linear bandits. *Advances in Neural Information*
367 *Processing Systems*, 37:37791–37822, 2024.
- 368 Eray Can Elumar, Cem Tekin, and Osman Yağın. Multi-armed bandits with costly probes.
369 *IEEE Transactions on Information Theory*, 71(1):618–643, 2024.
- 370 Diego Feijer, Himan Abdollahpouri, Sanket Gupta, Alexander Clare, Yuxiao Wen, Todd
371 Wasson, Maria Dimakopoulou, Zahra Nazari, Kyle Kretschman, and Mounia Lalmas.
372 Calibrated recommendations with contextual bandits. *arXiv preprint arXiv:2509.05460*,
373 2025.
- 374 Yunlong Hou, Vincent Y Tan, and Zixin Zhong. Almost minimax optimal best arm identifi-
375 cation in piecewise stationary linear bandits. *Advances in Neural Information Processing*
376 *Systems*, 37:128967–129041, 2024.
- 377 Kyoungseok Jang, Chicheng Zhang, and Kwang-Sung Jun. Efficient low-rank matrix esti-
378 mation, experimental design, and arm-set-dependent low-rank bandits. In *Proceedings of*
379 *the 41st International Conference on Machine Learning*, pages 21329–21372, 2024.
- 380 Yassir Jedra, William Réveillard, Stefan Stojanovic, and Alexandre Proutiere. Low-rank
381 bandits via tight two-to-infinity singular subspace recovery. In *International Conference*
382 *on Machine Learning*, pages 21430–21485. PMLR, 2024.
- 383 Kwang-Sung Jun, Rebecca Willett, Stephen Wright, and Robert Nowak. Bilinear bandits
384 with low-rank structure. In *International Conference on Machine Learning*, pages 3163–
385 3172. PMLR, 2019.
- 386 Yue Kang, Cho-Jui Hsieh, and Thomas Chun Man Lee. Efficient frameworks for generalized
387 low-rank matrix bandit problems. *Advances in Neural Information Processing Systems*,
388 35:19971–19983, 2022.
- 389 Yeoneung Kim, Insoon Yang, and Kwang-Sung Jun. Improved regret analysis for variance-
390 adaptive linear bandits and horizon-free linear mixture mdps. *Advances in Neural Infor-*
391 *mation Processing Systems*, 35:1060–1072, 2022.
- 392 Junpei Komiyama, Edouard Fouché, and Junya Honda. Finite-time analysis of globally
393 nonstationary multi-armed bandits. *Journal of Machine Learning Research*, 25(112):1–
394 56, 2024.
- 395 Beatrice Laurent and Pascal Massart. Adaptive estimation of a quadratic functional by
396 model selection. *Annals of statistics*, pages 1302–1338, 2000.
- 397 Pan Liu, Yaguang Li, and Jialiang Li. Change surface regression for nonlinear subgroup
398 identification with application to warfarin pharmacogenomics data. *Biometrics*, 81(1):
399 ujae169, 2025.
- 400 Yangyi Lu, Amirhossein Meisami, and Ambuj Tewari. Low-rank generalized linear bandit
401 problems. In *International conference on artificial intelligence and statistics*, pages 460–
402 468. PMLR, 2021.

- 403 Yoan Russac, Claire Vernade, and Olivier Cappé. Weighted linear bandits for non-stationary
404 environments. *Advances in Neural Information Processing Systems*, 32, 2019.
- 405 Michael J Rybak, Jennifer Le, Thomas P Lodise, Donald P Levine, John S Bradley, Catherine
406 Liu, Bruce A Mueller, Manjunath P Pai, Annie Wong-Beringer, John C Rotschafer,
407 et al. Therapeutic monitoring of vancomycin for serious methicillin-resistant staphylococ-
408 cus aureus infections: a revised consensus guideline and review by the american society
409 of health-system pharmacists, the infectious diseases society of america, the pediatric
410 infectious diseases society, and the society of infectious diseases pharmacists. *American
411 journal of health-system pharmacy*, 77(11):835–864, 2020.
- 412 Yuta Saito, Shunsuke Aihara, Megumi Matsutani, and Yusuke Narita. Open bandit dataset
413 and pipeline: Towards realistic and reproducible off-policy evaluation. In *Thirty-fifth
414 Conference on Neural Information Processing Systems Datasets and Benchmarks Track
415 (Round 2)*, 2021.
- 416 Yevgeny Seldin, Peter Bartlett, Koby Crammer, and Yasin Abbasi-Yadkori. Prediction
417 with limited advice and multiarmed bandits with paid observations. In *International
418 Conference on Machine Learning*, pages 280–287. PMLR, 2014.
- 419 Stefan Stojanovic, Yassir Jedra, and Alexandre Proutiere. Spectral entry-wise matrix esti-
420 mation for low-rank reinforcement learning. *Advances in Neural Information Processing
421 Systems*, 36:77056–77070, 2023.
- 422 Joel A Tropp. Freedman’s inequality for matrix martingales. *Electronic Communications
423 in Probability*, 16:262–270, 2011.
- 424 Aaron D Tucker, Caleb Biddulph, Claire Wang, and Thorsten Joachims. Bandits with costly
425 reward observations. In *Uncertainty in Artificial Intelligence*, pages 2147–2156. PMLR,
426 2023.

427 Appendix: Table of Contents

428	A. Comparison with prior work	p. 12
429	B. Proofs	p. 12
430	B.1 Preliminaries and notation.....	p. 14
431	B.2 Concentration of the lifted probe moment.....	p. 15
432	B.3 Necessity of the probe-side conditions.....	p. 17
433	B.4 Lemmas for the interleaved analysis.....	p. 18
434	B.5 Proof of Theorem 4.1 (costed dynamic regret of Alg. 1).....	p. 19
435	B.6 Proof of Corollary 4.4 (rank-adaptive).....	p. 20
436	B.7 Proof of Proposition 4.5 (SPSC-Adaptive).....	p. 20
437	C. Probe distributions: scaled sphere and Gaussian truncation	p. 21
438	D. Algorithms: pseudocode and adaptive variant	p. 21
439	E. Experimental setup	p. 23
440	F. Per-cell tables: every dataset, every cell	p. 24
441	G. Sensitivity and robustness studies	p. 28
442	H. Russac benchmark	p. 31
443	I. Warfarin: random-subspace ablation	p. 31
444	J. BOSS/Jedra adaptation	p. 33
445	K. Vancomycin: full per-rank tables	p. 33
446	L. When the adaptive variant wins	p. 33

447 A Comparison with prior work

448 This appendix positions SPSC against published work in low-rank bandits, nonstationary
449 bandits, and active-query bandits via three tables: Table 5 (head-to-head with seven recent
450 works on regret form, change-point adaptation, and probing budget), Table 6 (compact po-
451 sitioning of SPSC’s combined setting against the three closest *problem classes*), and Table 7
452 (regret-rate comparison with five published baselines that appear in our experiments). The
453 take-away is that the existing literature splits along three axes (low-rank, nonstationary,
454 costly probing) and no prior method addresses all three; SPSC is the first to do so with
455 a self-contained identification chain that turns rank-one scalar rewards into matrix-valued
456 evidence.

457 B Proofs

458 This section proves the statements of §4 in dependency order:

- 459 1. §B.1 fixes notation and establishes the probe-time excitation lemma that all proofs de-
460 pend on.
- 461 2. §B.2 derives the load-bearing *identification chain*: closed-form inverse of the probe oper-
462 ator $\mathcal{K}$, quadratic-measurement identity, lifted probe sample’s near-unbiasedness, matrix
463 Bernstein/Freedman, and the projector error via Davis–Kahan.
- 464 3. §B.4 states and proves the interleaved-analysis lemmas: projected nonstationary UCB
465 with time-varying basis (Lem. B.16), uniform prefix subspace concentration (Lem. B.17),
466 and cumulative interleaved subspace error (Lem. B.18).

Table 5: Comparison with related literature. “Prob.” indicates a probabilistic regret bound; “CP?” indicates change-point adaptation.

Reference	Venue/ Year	Regret / Theory	Prob.	CP?	Difference vs. ours
Kang et al. [2022]	NeurIPS’22	$\tilde{O}\left(\sqrt{(d_1 + d_2)MrT}\right)$	✓	✗	Stationary; probing not budgeted.
Stojanovic et al. [2023]	NeurIPS’23	Entry-wise / subspace recovery	✗/✓	✗	Estimation focus; no probe budget.
Jedra et al. [2024]	ICML’24	$\tilde{O}\left(r^{5/4}(m+n)^{3/4}\sqrt{T}\right)$	✓	✗	Stationary; no re-learning.
Jang et al. [2024]	ICML’24	Arm-set-dependent regret	✓	✗	Stationary; no CP detection.
Abbasi-Yadkori et al. [2023]	JMLR’23	$\tilde{O}(\sqrt{KN(S+1)})$	✗	✓	Unstructured K -arm; no low-rank.
Komiyama et al. [2024]	JMLR’24	ADR-bandit finite-time	✗	✓	Unstructured; no low-rank + probing.
Hou et al. [2024]	NeurIPS’24	Fixed-confidence sample complexity	✓	✓	BAI objective; no low-rank.

Table 6: Compact positioning of SPSC’s combined setting.

Work class	Low-rank	Nonstationary	Costly probing	Main gap vs. ours
Low-rank bandits	✓	✗	✗	Fixed subspace; no piecewise changes; no sensing-cost budget.
Nonstationary linear/contextual bandits	✗	✓	✗	Adapt in ambient parameter space; ignore latent low-rank structure.
Costly-information / active-query bandits	✗	✗	✓	Price information acquisition; do not address subspace recovery under changing representations.
This paper	✓	✓	✓	Combines low-rank structure, piecewise-stationary adaptation, and explicit costly probing in a single identification-to-exploitation framework.

Table 7: Regret-scaling comparison with the published baselines that appear in our experiments.

Algorithm	Reference	Regret scaling	Key difference vs. SPSC
LinUCB	Abbasi-Yadkori et al. [2011]	$O(d\sqrt{T})$	Ambient; no subspace exploitation.
LowOFUL	Jun et al. [2019]	$\tilde{O}(r\sqrt{T})$	Needs oracle subspace B .
LowRank-Reward	Lu et al. [2021]	$\tilde{O}(r\sqrt{T})$	Stationary; learns B from rewards.
BOSS	Duong et al. [2024]	$\tilde{O}(r\sqrt{dT})$	Stationary; no online change adaptation.
Jedra	Jedra et al. [2024]	$\tilde{O}(r^{5/4}(m+n)^{3/4}\sqrt{T})$	Stationary; no probe cost.
SPSC (ours)	This work	$\tilde{O}(r\sqrt{T} + T^{2/3})$	Learns a <i>changing</i> B via priced probes.

467 4. §B.5 proves the main costed regret bound (Thm. 4.1) via a five-step decomposition: probe
 468 regret, exploitation regret via Lem. B.16, per-segment bound, probe-estimation tradeoff
 469 with burn-in, and sum over segments.

470 5. §B.6 derives the rank-adaptive corollary (Cor. 4.4) via Weyl's inequality on the matrix-
 471 Bernstein bound.

472 6. §B.7 lifts the guarantee to SPSC-Adaptive (Prop. 4.5) by controlling the CUSUM detec-
 473 tor.

474 The windowed projected self-normalized bound (Lem. B.3), the windowed projected poten-
 475 tial condition (Assumption B.2), and the projected nonstationary UCB lemma (Lem. B.16)
 476 are the linear-bandit ingredients; constants that do not affect leading order are absorbed
 477 into $\tilde{\mathcal{O}}(\cdot)$.

478 B.1 Preliminaries and notation

479 Throughout, $\mathcal{I}_k = [\tau_{k-1}, \tau_k)$ denotes segment k of length $\ell_k := |\mathcal{I}_k|$, so $\sum_k \ell_k = T$. Inside
 480 $\mathcal{I}_k$ the learner performs $m_k = |\mathcal{T}_k|$ probe rounds ($\mathcal{T}_k := \mathcal{T}_{\text{probe}} \cap \mathcal{I}_k$) and $n_k = \ell_k - m_k$
 481 exploitation rounds ($E_k := \mathcal{I}_k \setminus \mathcal{T}_k$). We write $\hat{P}_t = \hat{U}_t \hat{U}_t^\top$ and $P_k^* = B_k^* B_k^{*\top}$ for the estimated
 482 and true rank- r projectors, and $\varepsilon_k := \|\hat{P}_k - P_k^*\|_{\text{op}}$. We work on the high-probability event
 483 $\mathcal{G}_w := \{\max_{t \leq T} \|w_t\|_2 \leq S_w\}$. For sub-Gaussian innovations one may take $S_w = \tilde{\mathcal{O}}(\sqrt{r})$
 484 via union bound (the failure probability is absorbed into the overall δ budget); for bounded
 485 innovations $\mathcal{G}_w$ holds deterministically. In what follows we implicitly condition on $\mathcal{G}_w$.

486 The *predictable second moment* is $\tilde{M}_t := \mathbb{E}[\theta_t \theta_t^\top \mid \mathcal{H}_{t-1}] = B_k^* \mathbb{E}[w_t w_t^\top \mid \mathcal{H}_{t-1}] (B_k^*)^\top$, and
 487 the probe-time average is $\bar{M}_k^{\text{probe}} := m_k^{-1} \sum_{t \in \mathcal{T}_k} \tilde{M}_t$. By the non-degenerate-innovation
 488 assumption $\Sigma_{\eta,k} \succ 0$ of §2, $\bar{M}_k^{\text{probe}}$ has rank r and its r -th eigenvalue $\lambda_{\min} > 0$ is uniformly
 489 bounded below in k (Lem. B.1 below; stability of A_k is used only to bound $\|w_t\| \leq S_w$, not
 490 for this lower bound). Constants $C, C', \dots$ depend only on $\sigma_\varepsilon, R_{\mathcal{A}}, L, \rho_Q, S_w, \lambda_{\min}$ and may
 491 change line to line.

492 **Lemma B.1** (Predictable probe-time excitation). *For each segment k ,*

$$\lambda_r \left(\frac{1}{m_k} \sum_{t \in \mathcal{T}_k} \mathbb{E}[w_t w_t^\top \mid \mathcal{H}_{t-1}] \right) \geq \lambda_{\min} > 0 \quad \text{almost surely.}$$

493 *Proof.* For each $t \in \mathcal{I}_k$, the conditional second moment satisfies $\mathbb{E}[w_t w_t^\top \mid \mathcal{H}_{t-1}] =$
 494 $A_k w_{t-1} w_{t-1}^\top A_k^\top + \Sigma_{\eta,k} \succeq \Sigma_{\eta,k}$ (LDS recursion + non-degenerate innovations). Hence
 495 $\lambda_r(\mathbb{E}[w_t w_t^\top \mid \mathcal{H}_{t-1}]) \geq \lambda_r(\Sigma_{\eta,k}) =: \lambda_{\min} > 0$ uniformly across rounds. Averaging over any
 496 probe subset $\mathcal{T}_k \subseteq \mathcal{I}_k$ preserves this PSD lower bound: $\lambda_r(m_k^{-1} \sum_{t \in \mathcal{T}_k} \mathbb{E}[w_t w_t^\top \mid \mathcal{H}_{t-1}]) \geq$
 497 $\lambda_{\min}$ almost surely. Stability of A_k is therefore not needed for the lemma; the innovation
 498 covariance alone suffices. $\square$

499 **Notation for the interleaved analysis.** For each segment $\mathcal{I}_k = [\tau_{k-1}, \tau_k)$, recall $\ell_k =$
 500 $|\mathcal{I}_k|$, $\mathcal{T}_k = \mathcal{T}_{\text{probe}} \cap \mathcal{I}_k$, $m_k = |\mathcal{T}_k|$, $E_k = \mathcal{I}_k \setminus \mathcal{T}_k$, $n_k = |E_k| = \ell_k - m_k$. For $t \in \mathcal{I}_k$,
 501 let $q_k(t) := |\mathcal{T}_k \cap [\tau_{k-1}, t)|$ be the number of probes collected in segment k strictly before
 502 round t . Denote $P_k^* = B_k^* B_k^{*\top}$, $\hat{P}_t = \hat{U}_t \hat{U}_t^\top$, and $\varepsilon_{k,t} := \|\hat{P}_t - P_k^*\|_{\text{op}}$. The within-segment
 503 variation is $V_k^{\text{in}} := \sum_{s, s+1 \in \mathcal{I}_k} \|\theta_{s+1} - \theta_s\|_2$, $V_{\text{in}} := \sum_k V_k^{\text{in}}$. Let $R_Q := \sup_{u \in \text{supp}(Q)} \|u\|_2$
 504 ($R_Q = \sqrt{d}$ for scaled-sphere probes; $R_Q \leq R_{\mathcal{A}}$ if probes are feasible actions covered by the
 505 action-radius bound).

506 **Balanced interleaved probe schedule.** The probe schedule is balanced: there exists
 507 an absolute constant a_0 such that, for every segment k and every $j \in \{0, 1, \dots, m_k\}$,

$$|\{t \in E_k : q_k(t) = j\}| \leq a_0 \lceil \ell_k / m_k \rceil. \quad (\text{B})$$

508 Uniformly spaced probes satisfy this condition.

Windowed projected potential. For each exploit round $t \in E_k$, Alg. 1 uses the *current* basis $\widehat{U}_{t-1}$ and re-projects the windowed history through this same basis: $z_t(x) := \widehat{U}_{t-1}^\top x$, $\widetilde{V}_t := \lambda I_r + \sum_{s \in \mathcal{W}_t} z_t(x_s) z_t(x_s)^\top$, $w_t(x) := \|z_t(x)\|_{\widetilde{V}_t^{-1}}$.

Assumption B.2 (Windowed projected potential). *For each segment $\mathcal{I}_k$,*

$$\sum_{t \in E_k} \min\{1, w_t(x_t)^2\} \leq C_{\text{pot}} (n_k/W + 1) r \log\left(1 + \frac{WR_A^2}{\lambda r}\right), \quad (6)$$

with the convention that the factor $(n_k/W + 1)$ is interpreted as 1 and W is replaced by n_k when $W \geq n_k$. For fixed basis this is the standard elliptical-potential lemma [Abbasi-Yadkori et al., 2011, Lem. 11] (resp. its sliding-window form when $W < n_k$); the moving-basis case reduces to controlling the cumulative log-determinant change under basis updates, which follows from prefix concentration (Lem. B.17) via a routine Lipschitz bound.

Lemma B.3 (Windowed projected self-normalized concentration). *For a fixed $a^* \in \mathbb{R}^r$ and σ_ε -sub-Gaussian noise on the projected covariates $z_t = \widehat{U}_t^\top x_t \in \mathbb{R}^r$, the windowed ridge estimator $\widehat{a}_t = \widetilde{V}_t^{-1} \widetilde{b}_t$ satisfies, with probability at least $1 - \delta/K$ per segment k and simultaneously for all $t \in E_k$,*

$$\|\widehat{a}_t - a^*\|_{\widetilde{V}_t} \leq \beta_t^{(r,W)}, \quad \beta_t^{(r,W)} := \sigma_\varepsilon \sqrt{r \log(1 + WR_A^2/(\lambda r)) + 2 \log(2K/\delta)} + \sqrt{\lambda} S_w.$$

This is the windowed projected specialization of Abbasi-Yadkori et al. [2011, Thm. 1, 2].

B.2 Concentration of the lifted probe moment (load-bearing step)

This is the non-standard piece of the analysis; we derive it in full. We work with scaled-sphere probes $u_t = \sqrt{d} v_t$, $v_t \sim \text{Unif}(\mathbb{S}^{d-1})$, which satisfy $\|u_t\|_2 = \sqrt{d}$ *deterministically* and $\mathbb{E}[u_t u_t^\top] = I_d$. The probe moment operator $\mathcal{K} : M \mapsto \mathbb{E}[(u^\top M u) u u^\top]$ is a linear bijection on the space of symmetric $d \times d$ matrices.

Lemma B.4 (Closed-form inverse of $\mathcal{K}$ on symmetric matrices). *For scaled-sphere probes $u = \sqrt{d} v$, $v \sim \text{Unif}(\mathbb{S}^{d-1})$, $\mathcal{K}$ acts on symmetric $d \times d$ matrices as $\mathcal{K}(M) = \frac{d}{d+2} (\text{tr}(M) I_d + 2M)$ with inverse $\mathcal{K}^{-1}(N) = \frac{d+2}{2d} N - \frac{\text{tr}(N)}{2d} I_d$. Moreover $\|\mathcal{K}^{-1}\|_{\text{op} \rightarrow \text{op}} \leq 1$ on symmetric inputs.*

Proof. Use the fourth-moment identity for the uniform sphere $\mathbb{E}[v_i v_j v_k v_l] = (\delta_{ij} \delta_{kl} + \delta_{ik} \delta_{jl} + \delta_{il} \delta_{jk}) / (d(d+2))$ and the rescaling $u = \sqrt{d} v$: $\mathbb{E}[(u^\top M u) u_k u_l] = d^2 \mathbb{E}[(v^\top M v) v_k v_l] = d^2 \cdot (2M_{kl} + \text{tr}(M) \delta_{kl}) / (d(d+2)) = \frac{d}{d+2} (2M_{kl} + \text{tr}(M) \delta_{kl})$ (using M symmetric). So $\mathcal{K}(M) = \frac{d}{d+2} (\text{tr}(M) I_d + 2M)$. Taking traces of $N = \mathcal{K}(M)$ gives $\text{tr}(N) = d \text{tr}(M)$, so $\text{tr}(M) = \text{tr}(N)/d$ and $M = \frac{d+2}{2d} N - \frac{\text{tr}(N)}{2d} I_d$.

For the operator-norm bound: $\mathcal{K}^{-1}(N)$ is diagonalized in the eigenbasis of N ; if N has eigenvalues $\mu_1, \dots, \mu_d$, then $\mathcal{K}^{-1}(N)$ has eigenvalues $\nu_i = \frac{d+2}{2d} \mu_i - \frac{\text{tr}(N)}{2d}$. Maximizing $|\nu_i|$ over $\|N\|_{\text{op}} = 1$ (i.e. $\max_i |\mu_i| = 1$): at $\mu_1 = 1, \mu_2 = \dots = \mu_d = -1$, $\text{tr}(N) = 2 - d$ and $\nu_1 = \frac{d+2-(2-d)}{2d} = 1$; at $\mu_1 = -1, \mu_2 = \dots = \mu_d = 1$, $\nu_1 = -1$. All other eigenvalue patterns give $|\nu_i| \leq 1$ by direct case analysis. So $\|\mathcal{K}^{-1}\|_{\text{op} \rightarrow \text{op}} = 1$ on symmetric matrices for $d \geq 2$. $\square$

Lemma B.5 (Quadratic measurement identity). *On every probe round $t \in \mathcal{T}_{\text{probe}}$, with $\delta_\sigma := \widehat{\sigma}^2 - \sigma_\varepsilon^2$ and $m_t := \mathbb{E}[\varepsilon_t \theta_t \mid \mathcal{H}_{t-1}, u_t]$,*

$$\mathbb{E}[s_t \mid \sigma(\mathcal{H}_{t-1}, u_t)] = u_t^\top \widetilde{M}_t u_t + 2u_t^\top m_t - \delta_\sigma, \quad \|m_t\| \leq \epsilon_\times.$$

Under exact probe conditions ($\delta_\sigma = \epsilon_\times = 0$) this collapses to $\mathbb{E}[s_t \mid \cdot] = u_t^\top \widetilde{M}_t u_t$.

Proof. Expand $y_t^2 = (u_t^\top \theta_t)^2 + 2\varepsilon_t u_t^\top \theta_t + \varepsilon_t^2$ and condition on $\sigma(\mathcal{H}_{t-1}, u_t)$; the middle term contributes $2u_t^\top m_t$ with m_t as defined in the lemma, bounded by $\epsilon_\times$ via the bounded state-noise coupling assumption (§2); the last term contributes σ_ε^2 , and subtracting $\widehat{\sigma}^2$ leaves $-\delta_\sigma$. Independence of u_t from $(\theta_t, \mathcal{H}_{t-1})$ gives $\mathbb{E}[(u_t^\top \theta_t)^2 \mid \mathcal{H}_{t-1}, u_t] = u_t^\top \widetilde{M}_t u_t$. $\square$

Lemma B.6 (Lifted probe sample is near-unbiased). *Let $G_t := \mathcal{K}^{-1}(s_t u_t u_t^\top)$ for scaled-sphere $u_t = \sqrt{d} v_t$, $v_t \sim \text{Unif}(\mathbb{S}^{d-1})$. Then $\mathbb{E}[G_t \mid \mathcal{H}_{t-1}] = \widetilde{M}_t + \widetilde{B}_t$ with bias*

$$\widetilde{B}_t = -\frac{\delta_\sigma}{d} I_d, \quad \|\widetilde{B}_t\|_{\text{op}} = \frac{|\delta_\sigma|}{d}.$$

The bias is a scaled identity, hence shifts every eigenvalue uniformly and does not rotate eigenvectors of $\widetilde{M}_t$.

Proof. By linearity of $\mathcal{K}^{-1}$ and Lemma B.5, $\mathbb{E}[G_t \mid \mathcal{H}_{t-1}] = \mathcal{K}^{-1}(\mathcal{K}(\widetilde{M}_t)) + 2\mathcal{K}^{-1}(\mathbb{E}[(u_t^\top m_t) u_t u_t^\top \mid \mathcal{H}_{t-1}]) - \delta_\sigma \mathcal{K}^{-1}(\mathbb{E}[u_t u_t^\top]) = \widetilde{M}_t + \widetilde{B}_t$, with $\widetilde{M}_t = \mathcal{K}^{-1}\mathcal{K}(\widetilde{M}_t)$.

Variance-bias term. $\mathbb{E}[u_t u_t^\top] = I_d$ exactly for scaled-sphere probes, so $-\delta_\sigma \mathcal{K}^{-1}(I_d) = -\delta_\sigma [\frac{d+2}{2d} I_d - \frac{d}{2d} I_d] = -\delta_\sigma \cdot \frac{2}{2d} I_d = -\frac{\delta_\sigma}{d} I_d$. This is a scaled identity, hence does not rotate eigenvectors of $\widetilde{M}_t$.

State-noise coupling term. The remaining contribution is $2\mathcal{K}^{-1}(\mathbb{E}[(u_t^\top m_t) u_t u_t^\top \mid \mathcal{H}_{t-1}])$. The (a, b) -entry of the inner expectation is $\sum_c (m_t)_c \mathbb{E}[(u_t)_a (u_t)_b (u_t)_c]$, a degree-three polynomial in u_t . For sphere-symmetric probes $u_t \mapsto -u_t$ leaves the distribution invariant, so every odd-degree moment vanishes and $\mathbb{E}[(u_t^\top m_t) u_t u_t^\top \mid \mathcal{H}_{t-1}] = 0$ exactly. Hence the coupling term contributes zero, and $\widetilde{B}_t = -\frac{\delta_\sigma}{d} I_d$. $\square$

Lemma B.7 (Uniform a.s. bound on G_t). *For scaled-sphere probes $\|u_t\| = \sqrt{d}$ deterministically; on the noise-truncation event $|\varepsilon_t| \leq L_\varepsilon := \sigma_\varepsilon \sqrt{2 \log(2T/\delta)}$ (which holds with probability $\geq 1 - \delta$ via union over T rounds), $|s_t| \leq R_s := (\sqrt{d} S_w + L_\varepsilon)^2 + \widehat{\sigma}^2$, and the centered lifted sample satisfies $\|\widetilde{X}_t\|_{\text{op}} \leq 2R_X$ with $R_X := d \cdot R_s + S_w^2$. So both the envelope and predictable variance are at most $2R_X$ and R_X^2 respectively.*

Proof. For scaled-sphere u_t , $\|u_t\| = \sqrt{d}$ exactly. $|s_t| = |(u_t^\top \theta_t + \varepsilon_t)^2 - \widehat{\sigma}^2| \leq (\sqrt{d} S_w + L_\varepsilon)^2 + \widehat{\sigma}^2 = R_s$ using $|u_t^\top \theta_t| \leq \|u_t\| \|\theta_t\| \leq \sqrt{d} S_w$ on $\|\theta_t\| \leq S_w$ (orthonormal B_k^* , $\|w_t\| \leq S_w$). Then $\|G_t\|_{\text{op}} = \|\mathcal{K}^{-1}(s_t u_t u_t^\top)\|_{\text{op}} \leq |s_t| \|u_t u_t^\top\|_{\text{op}} \cdot 1 = R_s \cdot d$ using $\|\mathcal{K}^{-1}\|_{\text{op} \rightarrow \text{op}} \leq 1$ and $\|u_t u_t^\top\|_{\text{op}} = d$. Adding the deterministic $S_w^2 = \|\widetilde{M}_t\|_{\text{op}}$ bound for the centering gives $\|\widetilde{X}_t\|_{\text{op}} \leq 2R_X$. $\square$

Theorem B.8 (Matrix Bernstein for $\widehat{M}_k$). *Let $\widehat{M}_k := m_k^{-1} \sum_{t \in \mathcal{T}_k} G_t$ and $X_t := G_t - \mathbb{E}[G_t \mid \mathcal{H}_{t-1}]$. Conditional on Lemma B.7, $\{X_t\}$ is a bounded matrix-valued MDS with $\|X_t\|_{\text{op}} \leq 2R_X$ a.s. and $\|\sum_{t \in \mathcal{T}_k} \mathbb{E}[X_t^2 \mid \mathcal{H}_{t-1}]\|_{\text{op}} \leq m_k R_X^2$. Here $\widetilde{B}$ denotes the segment-level bias; for scaled-sphere probes under the standing assumptions of §2, $\widetilde{B}_t$ from Lem. B.6 collapses to the deterministic constant $-(\delta_\sigma/d)I_d$, so we drop the subscript when no ambiguity arises and write $\widetilde{B}$ for the common value. Freedman's matrix inequality [Tropp, 2011] gives, with probability at least $1 - \delta$,*

$$\|\widehat{M}_k - \bar{M}_k^{\text{probe}} - \widetilde{B}\|_{\text{op}} \leq 2R_X \sqrt{\log(2d/\delta)/m_k} + \frac{2R_X \log(2d/\delta)}{3m_k}. \quad (7)$$

Proof. The MDS property is Lemma B.6 (shifted by the segment-level bias $\widetilde{B}$ from Lemma B.6, which is deterministic and thus enters as a constant offset). The a.s. bound $\|X_t\|_{\text{op}} \leq 2R_X$ follows from $\|G_t\|_{\text{op}} \leq R_X$ and $\|\mathbb{E}[G_t]\|_{\text{op}} \leq R_X$. The predictable variance bound follows from $\mathbb{E}[X_t^2] \preceq \mathbb{E}[G_t^2] \preceq R_X^2 I$. Applying the matrix Freedman inequality to the MDS $\{X_t\}$ in $\mathbb{R}_{\text{sym}}^{d \times d}$ with these two bounds yields (7); the leading $\sqrt{\log(2d/\delta)/m_k}$ rate dominates the $\log(2d/\delta)/m_k$ term for $m_k \gtrsim \log(d/\delta)$. $\square$

Proposition B.9 (Segment factorization). $\bar{M}_k^{\text{probe}} = B_k^* \bar{S}_k^{\text{probe}} (B_k^*)^\top$ where $\bar{S}_k^{\text{probe}} := m_k^{-1} \sum_{t \in \mathcal{T}_k} \mathbb{E}[w_t w_t^\top \mid \mathcal{H}_{t-1}]$. In particular, $\text{range}(\bar{M}_k^{\text{probe}}) = \text{range}(B_k^*)$.

Proof. $\widetilde{M}_t = B_k^* \mathbb{E}[w_t w_t^\top \mid \mathcal{H}_{t-1}] (B_k^*)^\top$ for $t \in \mathcal{I}_k$; average over $\mathcal{T}_k$ and pull B_k^* out. Range equality uses $\lambda_r(\bar{S}_k^{\text{probe}}) \geq \lambda_{\min} > 0$ (Lem. B.1). $\square$

591 **Corollary B.10** (Projector concentration). *If m_k is large enough that $\|\widehat{M}_k - \bar{M}_k^{\text{probe}}\|_{\text{op}} \leq$*
 592 $\lambda_{\min}/2$, *then with probability $\geq 1 - \delta$,*

$$\varepsilon_k := \|\widehat{P}_k - P_k^*\|_{\text{op}} \leq C_{\text{sub}} \sqrt{\log(2d/\delta)/m_k} + \Delta_\sigma, \quad C_{\text{sub}} := \frac{8R_X}{\lambda_{\min}}, \quad \Delta_\sigma := \frac{4|\delta_\sigma|/d}{\lambda_{\min}}. \quad (8)$$

593 *For scaled-sphere probes, the variance-misspecification component $|\delta_\sigma|/d$ corresponds to a*
 594 *scaled-identity bias (Lem. B.6) that is parallel to I_d and therefore preserves the eigenvectors*
 595 *of $\bar{M}_k^{\text{probe}}$; applying Davis–Kahan with the biased mean as reference removes the $|\delta_\sigma|/d$*
 596 *contribution to Δ_σ at population level (Cor. 4.2).*

597 *Proof.* By Lemma B.6 (scaled-sphere), $\mathbb{E}[\widehat{M}_k \mid \mathcal{H}] = \bar{M}_k^{\text{probe}} + \widetilde{B}$ with $\widetilde{B} = -(\delta_\sigma/d)I_d$ (a
 598 scaled identity), so $\|\widetilde{B}\|_{\text{op}} = |\delta_\sigma|/d =: b_\sigma$. Combining with Theorem B.8, $\|\widehat{M}_k - \bar{M}_k^{\text{probe}}\|_{\text{op}} \leq$
 599 $2R_X \sqrt{\log(2d/\delta)/m_k} + b_\sigma$. $\bar{M}_k^{\text{probe}}$ has rank exactly r with r -th eigenvalue $\geq \lambda_{\min}$ (Lem. B.1
 600 via Prop. B.9); Davis–Kahan [Bhatia, 2013, Thm. VII.3.1] gives $\|\widehat{P}_k - P_k^*\|_{\text{op}} \leq 4\|\widehat{M}_k -$
 601 $\bar{M}_k^{\text{probe}}\|_{\text{op}}/\lambda_{\min}$, yielding (8). $\square$

602 **Lemma B.11** (Variance estimator from probe residuals). *Split the first N probes of segment*
 603 *k in half: form $\widehat{M}_k^{(N/2)}$ from the first $N/2$ probes (matrix-Bernstein estimate), and define*
 604 *the residual estimator*

$$\widehat{\sigma}^2 := \text{med}_{t \in \mathcal{T}_k, N/2 < t \leq N} \{ y_t^2 - u_t^\top \widehat{M}_k^{(N/2)} u_t \}.$$

605 *On the event of Theorem B.8 for $\widehat{M}_k^{(N/2)}$, with probability at least $1 - \delta$,*

$$|\widehat{\sigma}^2 - \sigma_\varepsilon^2| \leq C \sigma_\varepsilon^2 \sqrt{\log(1/\delta)/N}.$$

606 *Proof sketch.* By Lemma B.5, on probe round t , $\mathbb{E}[y_t^2 - u_t^\top \widetilde{M}_t u_t \mid \mathcal{H}_{t-1}, u_t] = \sigma_\varepsilon^2 + O(\epsilon_\times)$.
 607 Substituting $\widehat{M}_k^{(N/2)}$ for $\widetilde{M}_t$ contributes $|u_t^\top (\widehat{M}_k^{(N/2)} - \widetilde{M}_t) u_t| = O(\sqrt{1/N})$ on the matrix-
 608 Bernstein event. The squared noise ε_t^2 is sub-exponential with mean σ_ε^2 ; the median of
 609 $N/2$ i.i.d. samples concentrates at rate $\sigma_\varepsilon^2 \sqrt{\log(1/\delta)/N}$ by standard concentration of the
 610 empirical median. $\square$

611 This chain (Lemmas B.4–B.7, Theorem B.8, Corollary B.10, Proposition B.9) gives a self-
 612 contained subspace-recovery guarantee for SPSC.

613 B.3 Necessity of the probe-side conditions

614 The probe-side conditions of §2 (known probe distribution Q with full-dimensional coverage,
 615 known noise variance, and bounded state-noise coupling $\|\mathbb{E}[\varepsilon_t \theta_t]\| \leq \epsilon_\times$) are each individu-
 616 ally necessary for identifiability of the predictable second moment $\widetilde{M}_t$ from single-play scalar
 617 rewards. The three propositions below construct two distinct conditional second moments
 618 generating identical observation laws when the named condition is dropped.

619 *Scope.* These propositions establish identifiability obstructions in the unrestricted-second-
 620 moment problem and serve as necessity checks for the probe-side modeling assumptions.
 621 They are not lower bounds inside the rank- r LDS class of §2: the counterexample second
 622 moments need not be rank- r or compatible with a stable LDS w_t .

623 **Proposition B.12** (Non-identification without known noise variance). *Fix t . Suppose*
 624 *$\epsilon_\times = 0$ (exact state-noise orthogonality) and full probe coverage hold, but the conditional*
 625 *noise variance $v_t(u) := \mathbb{E}[\varepsilon_t^2 \mid \mathcal{H}_{t-1}, u]$ is completely unknown. Then $\widetilde{M}_t$ is not identifiable*
 626 *from the observable $g_t(u) := \mathbb{E}[y_t^2 \mid \mathcal{H}_{t-1}, u]$.*

627 *Proof.* $g_t(u) = u^\top \widetilde{M}_t u + v_t(u)$. For any small $\delta \in (0, \sigma_\varepsilon^2/L^2]$, set $H := \delta I_d$ and define
 628 $\widetilde{M}'_t := \widetilde{M}_t + H$ and $v'_t(u) := v_t(u) - u^\top H u = v_t(u) - \delta \|u\|^2$. On the truncation support
 629 $\|u\| \leq L$, $v'_t(u) \geq \sigma_\varepsilon^2 - \delta L^2 \geq 0$, so v'_t is a valid conditional variance. The two pairs $(\widetilde{M}_t, v_t)$
 630 and $(\widetilde{M}'_t, v'_t)$ produce identical observables g_t , so $\widetilde{M}_t \neq \widetilde{M}'_t$ are indistinguishable. $\square$

Proposition B.13 (Non-identification without bounded state-noise coupling). *Fix t . Suppose σ_ε^2 is known and probe coverage is full, but the cross-moment $m_t(u) := \mathbb{E}[\theta_t \varepsilon_t \mid \mathcal{H}_{t-1}, u]$ is arbitrary (no bound $\epsilon_\times$). Then $\widetilde{M}_t$ is not identifiable from $h_t(u) := \mathbb{E}[y_t^2 - \sigma_\varepsilon^2 \mid \mathcal{H}_{t-1}, u]$.*

Proof. $h_t(u) = u^\top \widetilde{M}_t u + 2u^\top m_t(u)$. For any nonzero symmetric H , let $\widetilde{M}'_t := \widetilde{M}_t + H$ and $m'_t(u) := m_t(u) - \frac{1}{2}Hu$. Then $u^\top \widetilde{M}'_t u + 2u^\top m'_t(u) = u^\top \widetilde{M}_t u + u^\top H u + 2u^\top m_t(u) - u^\top H u = h_t(u)$ for every u , so the two pairs are observationally equivalent. $\square$

Proposition B.14 (Non-identification without full-dimensional coverage). *Suppose probe actions u_t lie in a proper subspace $S \subsetneq \mathbb{R}^d$ almost surely. Then $\widetilde{M}_t$ is not identifiable from either linear rewards y_t or quadratic observations y_t^2 .*

Proof. Pick any nonzero symmetric H supported on $S^\perp$ (i.e. $H = P_{S^\perp} H P_{S^\perp} \neq 0$, possible since $S^\perp \neq \{0\}$). For every $u \in S$, $Hu \in S^\perp$ and $u^\top Hu = 0$, so $u^\top (\widetilde{M}_t + H)u = u^\top \widetilde{M}_t u$ and $u^\top (\theta_t + Hz) = u^\top \theta_t$ for any $z \in S^\perp$. Thus the pairs $(\theta_t, \widetilde{M}_t)$ and $(\theta_t + Hz, \widetilde{M}_t + H)$ generate identical observations on all probe rounds, so $\widetilde{M}_t$ is not identifiable. $\square$

Remark B.15 (Joint necessity). *No two of the three conditions jointly suffice: Prop. B.12 fails only on (i); Prop. B.13 only on (ii); Prop. B.14 only on (iii). Each counterexample satisfies the remaining two conditions, so the three form a minimal identifiability set for the single-play quadratic measurement model.*

B.4 Lemmas for the interleaved analysis

Lemma B.16 (Projected nonstationary UCB with time-varying basis). *Fix a segment $\mathcal{I}_k$. For each exploit round $t \in E_k$, define $U_t := \widehat{U}_{t-1}$, $P_t := U_t U_t^\top$, $z_t(x) := U_t^\top x$, $a_t^* := U_t^\top \theta_t$, and (as in Alg. 1) $\widetilde{V}_t = \lambda I_r + \sum_{s \in \mathcal{W}_t} z_t(x_s) z_t(x_s)^\top$, $\widehat{a}_t = \widetilde{V}_t^{-1} \sum_{s \in \mathcal{W}_t} z_t(x_s) y_s$. Assume on segment $\mathcal{I}_k$: (i) the windowed projected self-normalized event of Lem. B.3 holds uniformly over E_k ; (ii) Assumption B.2 holds; (iii) $\|\theta_t\| \leq S_w$, $\|x\| \leq R_{\mathcal{A}}$ for $x \in \mathcal{A}_t$. Suppose Alg. 1 uses an optimism correction satisfying*

$$\gamma_t \geq C_\gamma S_w \left(1 + R_{\mathcal{A}} \sqrt{W/\lambda}\right) \varepsilon_{k,t-1} + C_\gamma R_{\mathcal{A}} V_{k,t}(W), \quad (9)$$

where $L_W := \log(1 + WR_{\mathcal{A}}^2/(\lambda r))$ and $V_{k,t}(W) := \sum_{s, s+1 \in \mathcal{I}_k, t-W \leq s < t} \|\theta_{s+1} - \theta_s\|_2$. Then on the same event,

$$\sum_{t \in E_k} \left(\max_{x \in \mathcal{A}_t} x^\top \theta_t - x_t^\top \theta_t \right) \leq \widetilde{\mathcal{O}}(r\sqrt{n_k}) + C_{\text{tv}} R_{\mathcal{A}} S_w \left(1 + R_{\mathcal{A}} \sqrt{W/\lambda}\right) \sum_{t \in E_k} \varepsilon_{k,t-1} + C R_{\mathcal{A}} W V_k^{\text{in}}. \quad (10)$$

Proof. Fix $t \in E_k$ and decompose, for $x \in \mathcal{A}_t$, $x^\top \theta_t = z_t(x)^\top a_t^* + \zeta_t(x)$ with $\zeta_t(x) = x^\top (I - P_t) \theta_t$. Since $\theta_t = P_k^* \theta_t$ within segment k and $\|(I - P_t) P_k^*\|_{\text{op}} = \|P_t - P_k^*\|_{\text{op}} = \varepsilon_{k,t-1}$ for two same-rank orthogonal projectors,

$$|\zeta_t(x)| \leq R_{\mathcal{A}} S_w \varepsilon_{k,t-1}. \quad (11)$$

Rewrite the windowed regression in basis U_t . For $s \in \mathcal{W}_t$, $y_s = z_t(x_s)^\top a_t^* + b_{s,t} + \varepsilon_s$ with bias $b_{s,t} = x_s^\top (I - P_t) \theta_s + x_s^\top P_t (\theta_s - \theta_t)$. The first term has $|x_s^\top (I - P_t) \theta_s| \leq R_{\mathcal{A}} S_w \varepsilon_{k,t-1}$ (again $\theta_s = P_k^* \theta_s$). The second is at most $R_{\mathcal{A}} \sum_{u=s}^{t-1} \|\theta_{u+1} - \theta_u\|_2$.

By definition, $\widehat{a}_t - a_t^* = \widetilde{V}_t^{-1} \sum_{s \in \mathcal{W}_t} z_t(x_s) \varepsilon_s + \widetilde{V}_t^{-1} \sum_{s \in \mathcal{W}_t} z_t(x_s) b_{s,t} - \lambda \widetilde{V}_t^{-1} a_t^*$. On the self-normalized event of Lem. B.3 the noise+ridge contribution has weighted norm $\leq \beta_t^{(r,W)} = \widetilde{\mathcal{O}}(\sqrt{r})$.

For the deterministic subspace-bias part, let $b_t^{\text{sub}} \in \mathbb{R}^{|\mathcal{W}_t|}$ have entries $x_s^\top (I - P_t) \theta_s$, with $\|b_t^{\text{sub}}\|_\infty \leq R_{\mathcal{A}} S_w \varepsilon_{k,t-1}$. Let Z_t be the matrix with rows $z_t(x_s)^\top$, $s \in \mathcal{W}_t$. Since

668 $\sum_{s \in \mathcal{W}_t} z_t(x_s) z_t(x_s)^\top = Z_t^\top Z_t \preceq \tilde{V}_t$, $\|\tilde{V}_t^{-1/2} Z_t^\top\|_{\text{op}} \leq 1$, hence

$$\left\| \sum_{s \in \mathcal{W}_t} z_t(x_s) b_{s,t}^{\text{sub}} \right\|_{\tilde{V}_t^{-1}}^2 = \|\tilde{V}_t^{-1/2} Z_t^\top b_t^{\text{sub}}\|_2^2 \leq \|b_t^{\text{sub}}\|_2^2 \leq |\mathcal{W}_t| \|b_t^{\text{sub}}\|_\infty^2 \leq W R_{\mathcal{A}}^2 S_w^2 \varepsilon_{k,t-1}^2,$$

669 so the corresponding $\tilde{V}_t^{-1}$ -norm is at most $R_{\mathcal{A}} S_w \sqrt{W} \varepsilon_{k,t-1}$. Therefore for any x ,
 670 $|z_t(x)^\top \tilde{V}_t^{-1} \sum_s z_t(x_s) b_{s,t}^{\text{sub}}| \leq R_{\mathcal{A}} S_w \sqrt{W} \varepsilon_{k,t-1} w_t(x)$, and since $w_t(x) \leq \|x\|/\sqrt{\lambda}$, this is
 671 bounded by $R_{\mathcal{A}} S_w \sqrt{W/\lambda} \varepsilon_{k,t-1} \|x\|$, a quantity dominated by the $\gamma_t \|x\|$ correction in (9).
 672 The drift-bias part contributes $\leq C R_{\mathcal{A}} W V_k^{\text{in}}$ in total by the standard windowed argument
 673 (each increment appears in $\leq W$ windows).

674 By the optimism index in Alg. 1, with $x_t^* = \arg \max_{x \in \mathcal{A}_t} x^\top \theta_t$, the instantaneous regret
 675 obeys

$$\Delta_t \leq 2\beta_t^{(r,W)} w_t(x_t) + C_{\text{tv}} R_{\mathcal{A}} S_w \left(1 + R_{\mathcal{A}} \sqrt{W/\lambda}\right) \varepsilon_{k,t-1} + d_t,$$

676 where $\sum_t d_t \leq C R_{\mathcal{A}} W V_k^{\text{in}}$. Cauchy-Schwarz together with Assumption B.2 gives
 677 $\sum_{t \in E_k} \beta_t^{(r,W)} w_t(x_t) \leq (\sup_{t \in E_k} \beta_t^{(r,W)}) \sqrt{n_k \cdot 2C_{\text{pot}} r L_W} = \tilde{\mathcal{O}}(r \sqrt{n_k})$ (using $\beta_t^{(r,W)} =$
 678 $\tilde{\mathcal{O}}(\sqrt{r})$). Summing yields (10). $\square$

679 **Lemma B.17** (Uniform prefix subspace concentration). *Under the probe-side assumptions*
 680 *of §2, scaled-sphere probes, exact centering $\hat{\sigma}^2 = \sigma_\varepsilon^2$, and $\lambda_{\min} > 0$, there exist constants*
 681 *$C_{\text{sub}}, C_0 > 0$ such that, with probability at least $1 - \delta/3$, simultaneously for all segments k*
 682 *and all rounds $t \in \mathcal{I}_k$,*

$$\varepsilon_{k,t} \leq 1 \text{ if } q_k(t) < q^*, \quad \varepsilon_{k,t} \leq C_{\text{sub}} \sqrt{\log(2KdT/\delta)/q_k(t)} \text{ if } q_k(t) \geq q^*,$$

683 where $q^* := C_0 \log(2KdT/\delta)/\lambda_{\min}^2$.

684 *Proof.* For a fixed segment k and probe-prefix size q , the lifted probe estimator
 685 $q^{-1} \sum_{s \in \mathcal{T}_k: s \leq t} G_s$ satisfies a matrix Bernstein/Freedman bound (Thm. B.8) at deviation
 686 $C \sqrt{\log(2KdT/\delta)/q}$ with probability $\geq 1 - \delta/(3KT)$. By Lem. B.1 the predictable probe-
 687 time r th eigenvalue is $\geq \lambda_{\min}$. Whenever $q \geq q^* = C_0 \log(2KdT/\delta)/\lambda_{\min}^2$ the pertur-
 688 bation is at most a fixed fraction of the eigengap, and Davis-Kahan (Cor. B.10) gives
 689 $\|\hat{P}_t - P_k^*\|_{\text{op}} \leq C_{\text{sub}} \sqrt{\log(2KdT/\delta)/q}$. A union bound over all K segments and all at most
 690 T probe prefixes yields the simultaneous statement; for $q < q^*$ we use the trivial bound
 691 $\|\hat{P}_t - P_k^*\|_{\text{op}} \leq 1$. $\square$

692 **Lemma B.18** (Cumulative interleaved subspace error). *On the event of Lemma B.17, for*
 693 *every segment k ,*

$$\sum_{t \in E_k} \varepsilon_{k,t-1} \leq C \sqrt{\log(2KdT/\delta)} \frac{\ell_k}{\sqrt{m_k}} + C \frac{q^* \ell_k}{m_k}. \quad (12)$$

694 *Proof.* Group exploit rounds by $j = q_k(t-1)$. By the balanced-schedule condition (B)
 695 each group has at most $a_0 \lceil \ell_k/m_k \rceil$ exploit rounds. For $j < q^*$ use $\varepsilon_{k,t-1} \leq 1$, contributing
 696 $C q^* \ell_k/m_k$. For $j \geq q^*$, Lemma B.17 gives $\varepsilon_{k,t-1} \leq C_{\text{sub}} \sqrt{\log(2KdT/\delta)/j}$, so $\sum_{j=q^*}^{m_k} j^{-1/2} \leq$
 697 $2\sqrt{m_k}$ yields the first term. $\square$

698 B.5 Proof of Theorem 4.1

699 Fix segment $\mathcal{I}_k$. We decompose the costed regret on $\mathcal{I}_k$ into probe regret and exploitation
 700 regret.

701 **Probe regret.** On a probe round $t \in \mathcal{T}_k$, Alg. 1 plays $x_t = u_t \sim Q$ rather than an action
 702 in $\mathcal{A}_t$. Since $\|x\| \leq R_{\mathcal{A}}$ for $x \in \mathcal{A}_t$, $\|u_t\| \leq R_Q$, and $\|\theta_t\| \leq S_w$, $\max_{x \in \mathcal{A}_t} x^\top \theta_t - u_t^\top \theta_t \leq$
 703 $(R_{\mathcal{A}} + R_Q) S_w$, plus probe cost c . So each probe contributes at most $A := (R_{\mathcal{A}} + R_Q) S_w + c$,
 704 giving $\text{Reg}_{\text{probe}}(\mathcal{I}_k) \leq A m_k$.

705 **Exploitation regret.** By Lemmas B.16 and B.18,

$$\text{Reg}_{\text{exploit}}(\mathcal{I}_k) \leq \tilde{\mathcal{O}}(r\sqrt{n_k}) + B \frac{\ell_k}{\sqrt{m_k}} + CR_{\mathcal{A}}S_w \left(1 + R_{\mathcal{A}}\sqrt{W/\lambda}\right) \frac{q^*\ell_k}{m_k} + CR_{\mathcal{A}}WV_k^{\text{in}},$$

706 absorbing constants and logs into $B = CR_{\mathcal{A}}S_w(1 + R_{\mathcal{A}}\sqrt{W/\lambda})\sqrt{\log(2KdT/\delta)}$.

707 **Per-segment bound.** Combining,

$$\text{DynReg}^{(c)}(\mathcal{I}_k) \leq \tilde{\mathcal{O}}(r\sqrt{n_k}) + Am_k + B \frac{\ell_k}{\sqrt{m_k}} + \tilde{\mathcal{O}}\left(\frac{q^*\ell_k}{m_k}\right) + CR_{\mathcal{A}}WV_k^{\text{in}}. \quad (13)$$

708 **Optimizing the probe-estimation tradeoff.** Ignoring the lower-order burn-in term,
 709 minimize $f_k(m) := Am + B\ell_k m^{-1/2}$. Setting $f'_k(m) = A - B\ell_k/(2m^{3/2}) = 0$ gives $m_k^* =$
 710 $(B\ell_k/(2A))^{2/3}$, and $Am_k^* + B\ell_k/\sqrt{m_k^*} \leq C A^{1/3} B^{2/3} \ell_k^{2/3}$. At m_k^* the burn-in contribution is
 711 $\tilde{\mathcal{O}}(\ell_k^{1/3}/\lambda_{\min}^2)$ (after absorbing A, B constants). If $\ell_k < q^*$, the segment is charged trivially
 712 by $\tilde{\mathcal{O}}(\ell_k)$, absorbed into the same burn-in order. Hence

$$\text{DynReg}^{(c)}(\mathcal{I}_k) \leq \tilde{\mathcal{O}}(r\sqrt{n_k}) + \tilde{\mathcal{O}}(A^{1/3} B^{2/3} \ell_k^{2/3}) + \tilde{\mathcal{O}}(\ell_k^{1/3}/\lambda_{\min}^2) + CR_{\mathcal{A}}WV_k^{\text{in}}. \quad (14)$$

713 **Sum over segments.** By Cauchy-Schwarz, $\sum_k r\sqrt{n_k} \leq r\sqrt{KT}$. By Hölder, $\sum_k \ell_k^{2/3} \leq$
 714 $K^{1/3} T^{2/3}$ and $\sum_k \ell_k^{1/3} \leq K^{2/3} T^{1/3}$, and $\sum_k V_k^{\text{in}} = V_{\text{in}}$. Summing (14) over k gives

$$\text{DynReg}_T^{(c)} \leq \tilde{\mathcal{O}}(r\sqrt{KT}) + \tilde{\mathcal{O}}(A^{1/3} B^{2/3} K^{1/3} T^{2/3}) + \tilde{\mathcal{O}}(K^{2/3} T^{1/3}/\lambda_{\min}^2) + O(R_{\mathcal{A}}WV_{\text{in}}).$$

715 Union-bounding the windowed self-normalized event, the prefix subspace concentration, and
 716 the noise event at total δ yields (5). For fixed K , the leading term is $\tilde{\mathcal{O}}(r\sqrt{T})$, the probe
 717 term is $\tilde{\mathcal{O}}(T^{2/3})$, and the burn-in term $\tilde{\mathcal{O}}(T^{1/3})$ is lower order. When $V_{\text{in}} = 0$, the third
 718 term vanishes and we recover Corollary 4.3. $\square$

719 B.6 Proof of Corollary 4.4

720 By Theorem B.8, with probability $\geq 1 - \delta$ every eigenvalue of $\widehat{M}_k$ is within $\tau_k^{\text{rank}} :=$
 721 $2R_X\sqrt{\log(2d/\delta)/m_k}$ of the corresponding eigenvalue of $\bar{M}_k^{\text{probe}} + \widehat{B}$ (Weyl's inequality ap-
 722 plied to the matrix in (7)). The population spectrum has exactly r eigenvalues $\geq \lambda_{\min}$ and
 723 the remaining $d - r$ are zero (Prop. B.9). Under the eigengap hypothesis $\lambda_{\min} \geq 4\tau_k^{\text{rank}}$,
 724 thresholding at $2\tau_k^{\text{rank}}$ returns exactly r indices with probability $\geq 1 - \delta$. Condition on this
 725 event; the regret analysis of Theorem 4.1 applies verbatim with the estimated rank equal to
 726 the true r , so (5) holds with an extra δ in the union bound. $\square$

727 B.7 Proof of Proposition 4.5 (SPSC-Adaptive)

728 The detector statistic of Alg. 2 compares two non-overlapping rolling-window lifted-moment
 729 estimators of length W_{det} , $\widehat{M}_t^{\text{recent}}$ and $\widehat{M}_t^{\text{past}}$, via $S_t := \|\widehat{M}_t^{\text{recent}} - \widehat{M}_t^{\text{past}}\|_{\text{op}}$, and triggers
 730 a reset whenever $S_t > 2\eta_{\text{det}}$, where $\eta_{\text{det}} := C_{\text{sub}}\sqrt{\log(dT/\delta_{\text{FA}})/W_{\text{det}}}$ is the calibration
 731 radius. The proof has two parts: bounding the false-alarm probability under H_0 (no change
 732 inside the comparison window) and bounding the detection delay under H_1 (a change of
 733 size $\Delta_k \geq 2b$ has occurred), and then propagating the resulting boundary error through
 734 Theorem 4.1.

735 **(i) False-alarm control under H_0 .** Fix a round t and condition on H_0 : the predictable
 736 second moment is constant across the past and recent windows that produced $\widehat{M}_t^{\text{past}}$ and
 737 $\widehat{M}_t^{\text{recent}}$. Apply Theorem B.8 (matrix Bernstein for the lifted probe estimator) to each
 738 window separately at confidence $\delta_{\text{FA}}/(2T)$ to get

$$\|\widehat{M}_t^{\text{recent}} - \bar{M}_t^{\text{recent}}\|_{\text{op}} \leq \eta_{\text{det}}, \quad \|\widehat{M}_t^{\text{past}} - \bar{M}_t^{\text{past}}\|_{\text{op}} \leq \eta_{\text{det}},$$

each with probability at least $1 - \delta_{\text{FA}}/(2T)$, where $\bar{M}_t^{\text{recent}}$ and $\bar{M}_t^{\text{past}}$ are the predictable averages on the corresponding windows. Under H_0 they coincide, so by the triangle inequality $S_t \leq \|\bar{M}_t^{\text{recent}} - \bar{M}_t^{\text{recent}}\|_{\text{op}} + \|\bar{M}_t^{\text{past}} - \bar{M}_t^{\text{past}}\|_{\text{op}} \leq 2\eta_{\text{det}}$ with probability at least $1 - \delta_{\text{FA}}/T$, i.e. $\Pr(S_t > 2\eta_{\text{det}} \mid H_0) \leq \delta_{\text{FA}}/T$. A union bound over $t \in [T]$ gives total false-alarm probability at most δ_{FA} .

(ii) Detection delay under H_1 . Let τ_k be a true change point with eigengap $\Delta_k := \|P_k^* - P_{k-1}^*\|_{\text{op}} \geq 2b$ for the threshold $b > 2\eta_{\text{det}}$. Take $t = \tau_k + W_{\text{det}}$: the recent window now lies entirely in segment k while the past window lies entirely in segment $k-1$. By the segment-factorization of the predictable probe moment (Proposition B.9) and probe-time excitation (Lemma B.1), $\|\bar{M}_t^{\text{recent}} - \bar{M}_t^{\text{past}}\|_{\text{op}} \geq \lambda_{\min} \Delta_k \geq 2b\lambda_{\min}$. By the same matrix-Bernstein concentration as in (i), $S_t \geq \|\bar{M}_t^{\text{recent}} - \bar{M}_t^{\text{past}}\|_{\text{op}} - 2\eta_{\text{det}} \geq 2b\lambda_{\min} - 2\eta_{\text{det}} > 2\eta_{\text{det}}$ with probability at least $1 - \delta_{\text{FA}}/T$ (using $b > 2\eta_{\text{det}}$ after absorbing $\lambda_{\min}$ into b). Therefore the detector triggers no later than $\hat{\tau}_k = \tau_k + W_{\text{det}}$, giving detection delay $D_k \leq W_{\text{det}}$.

(iii) Boundary-error propagation. On the detector good event $\mathcal{E}_{\text{det}} = \{\text{no false alarm and every change detected within delay } D_k \leq W_{\text{det}}\}$, which holds with probability at least $1 - \delta_{\text{FA}}$ by parts (i)-(ii), the estimated boundaries $\hat{\tau}_k$ satisfy $|\hat{\tau}_k - \tau_k| \leq W_{\text{det}}$. Splitting the horizon at $\hat{\tau}_k$ and applying Theorem 4.1 on each estimated segment, the only difference from the oracle-boundary analysis is the contribution of the rounds in $[\tau_k, \hat{\tau}_k)$ where the algorithm is still using the old basis $\hat{U}_{\tau_k-1}$ on data drawn from segment k . Each such round contributes at most $2R_{\mathcal{A}}S_w$ to the instantaneous regret, so $\sum_k 2R_{\mathcal{A}}S_w D_k \leq 2R_{\mathcal{A}}S_w K W_{\text{det}} = O(KW_{\text{det}})$. Combining,

$$\text{DynReg}_T^{(c)}(\text{Alg. 2}) \leq \text{DynReg}_T^{(c)}(\text{Alg. 1}) + O(KW_{\text{det}})$$

on $\mathcal{E}_{\text{det}}$, which is the claim. The total failure probability is $\delta + \delta_{\text{FA}}$, absorbed into the overall δ budget by re-tuning constants. $\square$

C Probe distributions: scaled sphere and Gaussian truncation

The analysis in §B.2 uses scaled-sphere probes $u_t = \sqrt{d}v_t$, $v_t \sim \text{Unif}(\mathbb{S}^{d-1})$, which satisfy $\|u_t\|_2 = \sqrt{d}$ exactly. No truncation is required: the matrix-Bernstein envelope $R_X = dR_s + S_w^2$ holds deterministically given the noise-truncation event of Lemma B.7.

In our implementation, probes are generated via $u \leftarrow \sqrt{d}\tilde{u}/\|\tilde{u}\|$ for $\tilde{u} \sim \mathcal{N}(0, I_d)$, which is exactly the scaled-sphere distribution.

For isotropic-Gaussian probes (no rescaling), $u_t \sim \mathcal{N}(0, I_d)$, the Laurent–Massart χ^2 tail [Laurent and Massart, 2000, Lem. 1] gives $\Pr(\|u_t\| > L) \leq \delta/T$ at $L = \sqrt{2d \log(2T/\delta)}$; union-bounding over $|\mathcal{T}_{\text{probe}}| \leq T$ gives $\Pr(\bigcap_t \{\|u_t\| \leq L\}) \geq 1 - \delta$. On this event the results of §B.2 apply with $\|u_t\| \leq L$ (replacing $\sqrt{d}$), adding only a $\log(T/\delta)$ factor to R_X and C_{sub} that is absorbed in $\tilde{\mathcal{O}}(\cdot)$.

D Algorithms: pseudocode and adaptive variant

This appendix gives the full pseudocode for SPSC (oracle boundaries, Alg. 1) and SPSC-Adaptive (online change-point detection, Alg. 2), and explains the adaptive design choices (detector form, threshold, recovery phase, burn-in) that tie the tuning back to the guarantee of Proposition 4.5.

778 D.1 SPSC (oracle boundaries)

Algorithm 1 SPSC: Single-Play Subspace-Calibrated Optimism (oracle boundaries)

```

1: Input: segments  $\{\mathcal{I}_k\}_{k=1}^K$ , probe schedule  $\mathcal{T}_{\text{probe}}$ ,  $\lambda > 0$ , probe distribution  $Q$ , rank  $r$ ,
   window  $W$ .
2: Initialize  $\hat{U}_0 \in \mathbb{R}^{d \times r}$  arbitrarily with orthonormal columns.
3: for  $k = 1, \dots, K$  do ▷ outer loop over segments
4:    $M_{\text{acc}} \leftarrow 0$ ;  $N_k \leftarrow 0$ ;  $\hat{U}_t \leftarrow \hat{U}_{t-1}$ . ▷ reset segment-local accumulator and probe
   counter
5:   for  $t \in \mathcal{I}_k$  do
6:     if  $t \in \mathcal{T}_{\text{probe}}$  then ▷ probe round
7:       draw  $u_t \sim Q$ ; play  $x_t = u_t$ ; observe  $y_t$ .
8:        $s_t \leftarrow y_t^2 - \hat{\sigma}^2$ ;  $G_t \leftarrow \mathcal{K}^{-1}(s_t u_t u_t^\top)$ .
9:        $M_{\text{acc}} \leftarrow M_{\text{acc}} + G_t$ ;  $N_k \leftarrow N_k + 1$ ;  $\hat{U}_t \leftarrow \text{TopEig}_r(M_{\text{acc}}/N_k)$ .
10:    else ▷ exploitation round
11:       $\mathcal{W}_t \leftarrow \{s \in \mathcal{I}_k \setminus \mathcal{T}_{\text{probe}} : t-W \leq s < t\}$ ;  $z_t(x) \leftarrow \hat{U}_{t-1}^\top x$ .
12:       $\tilde{V}_t \leftarrow \lambda I_r + \sum_{s \in \mathcal{W}_t} z_t(x_s) z_t(x_s)^\top$ ,  $\tilde{b}_t \leftarrow \sum_{s \in \mathcal{W}_t} z_t(x_s) y_s$ .
13:       $\hat{a}_t \leftarrow \tilde{V}_t^{-1} \tilde{b}_t$ ; choose  $x_t = \arg \max_{x \in \mathcal{A}_t} \{z_t(x)^\top \hat{a}_t + \beta_t^{(r,W)} \|z_t(x)\|_{\tilde{V}_t^{-1}} +$ 
         $\gamma_t \|x\|_2\}$ .
14:      observe  $y_t$ ;  $\hat{U}_t \leftarrow \hat{U}_{t-1}$ .
15:    end if
16:  end for
17: end for

```

779 D.2 SPSC-Adaptive (unknown boundaries)

780 SPSC-Adaptive recovers segment boundaries from probe data, removing the oracle assumption
 781 of Alg. 1.

782 **Detector.** At each probe round the algorithm maintains two non-overlapping rolling-
 783 window estimators of the lifted second moment: $\hat{M}_t^{\text{recent}}$ over the last n probes and $\hat{M}_t^{\text{past}}$
 784 over the segment-accumulated probes preceding that window. The statistic

$$S_t := \|\hat{M}_t^{\text{recent}} - \hat{M}_t^{\text{past}}\|_{\text{op}} \quad (15)$$

785 is compared to a threshold b ; whenever $S_t > b$ the algorithm declares a change and enters
 786 a RECOVERY phase.

787 **Recovery phase.** Recovery deploys m_{relearn} consecutive probe rounds with no exploitation.
 788 The block-probe design serves two purposes: (i) by Proposition B.9, m_{relearn}
 789 back-to-back probes produce a subspace estimate $\hat{U}$ with projector error $\|\hat{P} - P^\star\|_{\text{op}} =$
 790 $O(1/\sqrt{m_{\text{relearn}}})$ on a $1 - \delta$ event; (ii) by zeroing out stale buffers it prevents the windowed
 791 ridge estimator from carrying information from the previous regime into the new one. After
 792 m_{relearn} rounds the algorithm transitions back to NORMAL probe-exploitation interleaving
 793 with probe rate μ .

794 **Threshold calibration and detection delay.** Under H_0 (no change), matrix Freedman
 795 applied to the two windows gives $\Pr(S_t > 2\eta_{\text{det}}) \leq \delta_{\text{FA}}/T$ with $\eta_{\text{det}} = C_{\text{sub}} \sqrt{\log(dT/\delta_{\text{FA}})/n}$;
 796 choosing $b = 2\eta_{\text{det}}$ and union-bounding gives at most δ_{FA} false alarms across the horizon.
 797 Under H_1 at time τ_k , a separation $\Delta_k \geq 2b$ forces a fire within delay $D_{\text{max}} = 2n/\mu + \tau_{\text{burn}}$,
 798 where the $2n/\mu$ term is the time for the recent window to refill with post-change probes, and
 799 τ_{burn} is a short burn-in period after recovery that prevents the detector from firing on its
 800 own freshly-reset state. The overall regret overhead is $2R_{\mathcal{A}} S_w \sum_k D_k = O(KW_{\text{det}})$, which
 801 is lower-order in the rate of Proposition 4.5.

Algorithm 2 SPSC-Adaptive: unknown boundaries via CUSUM-style detection

```

1: Input: probe rate  $\mu \in (0, 1)$ , detection window  $n$ , threshold  $b$ , burn-in  $\tau_{\text{burn}}$ , relearning
   budget  $m_{\text{relearn}}$ ,  $W$ ,  $\lambda$ .
2: phase  $\leftarrow$  RECOVERY; rec  $\leftarrow$  0;  $\widehat{M}_{\text{acc}} \leftarrow 0$ .
3: for  $t = 1, \dots, T$  do
4:   if phase = RECOVERY then  $\triangleright$  block-probe to rebuild subspace
5:     probe: draw  $u_t$ , play  $x_t = u_t$ , observe  $y_t$ ; update  $\widehat{M}_{\text{acc}}$ ; rec  $+= 1$ .
6:     if rec =  $m_{\text{relearn}}$  then
7:        $\widehat{U} \leftarrow \text{TopEig}_r(\widehat{M}_{\text{acc}}/m_{\text{relearn}})$ ; reset buffers; phase  $\leftarrow$  NORMAL.
8:     end if
9:   else  $\triangleright$  interleave probes at rate  $\mu$ 
10:    if Bernoulli( $\mu$ ) probe round then
11:      update lifted estimator  $\widehat{M}_{\text{acc}}$  and refresh rolling windows.
12:      if  $S_t > b$  (Eq. 15) and detector is armed then
13:        phase  $\leftarrow$  RECOVERY; rec  $\leftarrow$  0; reset.
14:      end if
15:    else
16:      windowed projected ridge-UCB step as in Alg. 1 (exploit branch).
17:    end if
18:  end if
19: end for

```

802 **Tuning used in experiments.** Across all experiments reported in §5 and the appendix
 803 we use $\mu = 0.1$, $n = W_{\text{det}} = 50$, $\tau_{\text{burn}} = 100$, $m_{\text{relearn}} = 30$, and b set from the CUSUM
 804 threshold `cusum_threshold` = 3.0. These are held fixed across datasets and (d, r) cells; no
 805 per-dataset retuning is done. Sensitivity to the probe rate μ and the detection window n is
 806 reported in §G.

807 E Experimental setup

808 **Synthetic phase-transition benchmark.** $\theta_t = B_k^* w_t$ with piecewise-constant orthonormal
 809 $B_k^* \in \mathbb{R}^{d \times r}$. Per segment, latent state evolves as $w_t = A_k w_{t-1} + \eta_{t-1}$ with spectral radius
 810 $\rho(A_k) = 0.99$ and Gaussian innovation $\eta_{t-1} \sim \mathcal{N}(0, 0.04^2 I_r)$. Reward noise $\varepsilon_t \sim \mathcal{N}(0, 0.09)$;
 811 40 actions per round drawn i.i.d. uniformly on the unit sphere. $K=10$ segments of equal
 812 length over $T=5,000$. Probe distribution Q isotropic Gaussian, probe cost $c=0.1$, probe
 813 period 50.

814 **UCI datasets (Covertypes, Pendigits, Satimage).** Features are the raw attribute
 815 vectors $\phi(x) \in \mathbb{R}^{d_0}$ reduced to ambient dimension d via Gaussian random projection $R \in$
 816 $\mathbb{R}^{d \times d_0}$. For Covertypes, segments correspond to disjoint size- K cover-type subsets of the 7
 817 classes; for Pendigits and Satimage, to disjoint digit/class subsets, $K=10$. Actions at round
 818 t are $|\mathcal{A}_t| = 40$ examples drawn from the active segment’s class pool. Reward is centered
 819 target label. $\lambda=0.01$, probe period 10, window $W=400$, $\delta=0.05$, $c=0.02$.

820 **Pixelized MNIST / Fashion-MNIST.** Images flattened, pixel-normalized, then pro-
 821 jected to d via random projection. Segments are class-subset swaps as above. Same hyper-
 822 parameters as UCI datasets.

823 **MovieLens-100K with real ratings.** Standard MovieLens-100K user-item matrix cen-
 824 tered by per-item mean. Features $\phi(\text{user}, \text{item})$ are item-side latent factors from a rank- d
 825 SVD of the observed-entry mask; reward is the raw user rating (1–5), centered. Segments
 826 correspond to disjoint time-bucketed user cohorts, $K=10$. This is a *fully real* benchmark
 827 (we do *not* generate rewards from a low-rank surrogate), so the data is only approximately
 828 low-rank and provides a stringent stress test. $T=5,000$, same hyperparameters.

829 **Warfarin clinical dosing.** IWPC cohort statistics ($\sim 5,700$ patients); $d=93$ features (de-
830 mographics, VKORC1/CYP2C9 genotype, medications, interactions). Reward is nega-
831 tive absolute deviation from the therapeutic dose, $K=8$ patient-cohort segments, $T=5,000$,
832 $\lambda=1.0$, probe period 10, $W=200$, $c=0.1$, 10 seeds.

833 **Baseline hyperparameters.** All baselines inherit λ, δ from SPSC when applicable. D-
834 LinUCB uses discount factor 0.998; SW-LinUCB uses window W ; Restart-LinUCB restarts
835 every T/K rounds (i.e. oracle boundaries); LinTS/SW-LinTS use posterior variance $\sigma^2 = \delta^2$.
836 Oracle-LinUCB knows the true B_k^* and runs LinUCB in the projected $\mathbb{R}^r$.

837 **Subspace-aware baselines (LowOFUL, VOFUL): favorable adaptation.** Both
838 LowOFUL and VOFUL were designed for a *fixed* low-rank representation, so running them
839 in our piecewise setting requires a choice for the subspace they consume. To avoid conflating
840 subspace acquisition with exploitation, we give these baselines the same refreshed projec-
841 tor $\hat{U}$ that SPSC uses (recomputed every 20 rounds after a 30-round warm-up), so every
842 subspace-aware method sees an identical low-rank representation at every round. This is a
843 *favorable* adaptation, since it removes representation quality as a possible source of SPSC’s
844 empirical gains, and any remaining gap is attributable to SPSC’s windowed head and drift
845 handling rather than to better subspace estimation. Running LowOFUL/VOFUL with their
846 own stationary subspace estimators (not shown) yields strictly worse baseline numbers and
847 larger SPSC margins.

848 **Implementation and compute.** All experiments use NumPy/SciPy only. A full (d, r)
849 grid (9 cells, 10 seeds, nine per-cell baselines plus Oracle-LinUCB and the two SPSC variants;
850 BOSS and Jedra appear only in the §5.6 grid) takes 6–20 hours on a single 16-core CPU.
851 Coverttype is the longest ($T=10,000$).

852 F Per-cell tables: every dataset, every cell

853 Tables 8–15 give the full per-cell results (mean $\pm$ SE costed regret over 10 seeds) for every
854 method on every (d, r) cell. Cells where at least one SPSC variant beats every non-oracle
855 baseline are highlighted green; the best non-oracle baseline per cell is underlined.

Table 8: **Coverttype** per-cell results. Mean $\pm$ SE costed regret over 10 seeds. Green = SPSC variant beats every non-oracle baseline; underline = best non-oracle. $T=10,000$, $K=4$. *Reading note (applies to all per-cell tables in this appendix):* rank-insensitive methods (LinUCB, SW-LinUCB, D-LinUCB, Restart-LinUCB, LinTS, SW-LinTS) act in ambient space; their regret depends on d but not on r , hence bit-identical values across the rows of each d -block.

d	r	Oracle	SPSC	SPSC-Adp	LowOFUL	VOFUL	LR-Rew	SW-Lin	D-Lin	Rst-Lin	LinUCB	LinTS	SW-LinTS
5	1	2013 $\pm$ 138	2655 $\pm$ 85	2114 $\pm$ 44	3748 $\pm$ 162	3736 $\pm$ 164	2379 $\pm$ 136	2651 $\pm$ 110	2502 $\pm$ 98	2661 $\pm$ 124	2694 $\pm$ 126	3633 $\pm$ 159	3064 $\pm$ 113
	55	1	718 $\pm$ 51	1296 $\pm$ 21	1262 $\pm$ 32	1429 $\pm$ 32	1436 $\pm$ 32	1366 $\pm$ 44	1331 $\pm$ 20	1346 $\pm$ 21	1329 $\pm$ 21	1336 $\pm$ 21	1462 $\pm$ 25
10	1	2768 $\pm$ 29	4307 $\pm$ 62	4161 $\pm$ 63	5067 $\pm$ 83	5167 $\pm$ 130	5122 $\pm$ 61	4403 $\pm$ 41	4599 $\pm$ 57	4348 $\pm$ 37	4465 $\pm$ 49	5591 $\pm$ 78	5486 $\pm$ 73
	20	4004 $\pm$ 55	5769 $\pm$ 54	5786 $\pm$ 84	7120 $\pm$ 93	7001 $\pm$ 72	7538 $\pm$ 81	5703 $\pm$ 34	5971 $\pm$ 40	5669 $\pm$ 40	5800 $\pm$ 50	7825 $\pm$ 41	7581 $\pm$ 44
30	1	5202 $\pm$ 48	6961 $\pm$ 95	7051 $\pm$ 81	8890 $\pm$ 110	8770 $\pm$ 93	9527 $\pm$ 96	6653 $\pm$ 34	6928 $\pm$ 30	6594 $\pm$ 51	6743 $\pm$ 34	9615 $\pm$ 105	9345 $\pm$ 90
	105	1	554 $\pm$ 30	1023 $\pm$ 29	1010 $\pm$ 21	1134 $\pm$ 41	1124 $\pm$ 44	1026 $\pm$ 28	1075 $\pm$ 32	1082 $\pm$ 33	1073 $\pm$ 31	1075 $\pm$ 33	1109 $\pm$ 35
10	1	2154 $\pm$ 27	3452 $\pm$ 45	3405 $\pm$ 40	3775 $\pm$ 34	3758 $\pm$ 48	3869 $\pm$ 36	3920 $\pm$ 36	3959 $\pm$ 36	3909 $\pm$ 37	3943 $\pm$ 35	4198 $\pm$ 40	4168 $\pm$ 37
	20	3233 $\pm$ 36	4600 $\pm$ 62	4507 $\pm$ 46	5297 $\pm$ 67	5325 $\pm$ 68	5621 $\pm$ 43	5302 $\pm$ 40	5383 $\pm$ 32	5313 $\pm$ 30	5355 $\pm$ 35	5826 $\pm$ 42	5734 $\pm$ 42
30	1	4165 $\pm$ 23	5529 $\pm$ 71	5450 $\pm$ 67	6476 $\pm$ 33	6556 $\pm$ 51	6989 $\pm$ 54	6360 $\pm$ 38	6473 $\pm$ 33	6324 $\pm$ 42	6400 $\pm$ 37	7085 $\pm$ 48	7029 $\pm$ 46

Table 9: **Pendigits** per-cell results. $T=5,000$, $K=10$, 10 seeds.

d	r	Oracle	SPSC	SPSC-Adp	LowOFUL	VOFUL	LR-Rew	SW-Lin	D-Lin	Rst-Lin	LinUCB	LinTS	SW-LinTS
5	1	39.0 $\pm$ 2.0	2914 $\pm$ 159	5480 $\pm$ 391	9536 $\pm$ 199	9623 $\pm$ 138	6606 $\pm$ 77	1090 $\pm$ 9	1267 $\pm$ 8	1022 $\pm$ 14	1073 $\pm$ 15	6629 $\pm$ 42	4630 $\pm$ 41
	55	1	18.0 $\pm$ 1.0	4019 $\pm$ 79	4073 $\pm$ 200	5913 $\pm$ 46	5939 $\pm$ 64	4542 $\pm$ 48	3863 $\pm$ 11	4056 $\pm$ 9	3620 $\pm$ 7	3658 $\pm$ 9	4343 $\pm$ 16
10	1	951 $\pm$ 5	2774 $\pm$ 58	2682 $\pm$ 93	3703 $\pm$ 71	3655 $\pm$ 67	4246 $\pm$ 28	3863 $\pm$ 11	4056 $\pm$ 9	3620 $\pm$ 7	3658 $\pm$ 9	4343 $\pm$ 16	4648 $\pm$ 9
	20	1816 $\pm$ 4	2803 $\pm$ 60	2490 $\pm$ 56	3689 $\pm$ 54	3854 $\pm$ 58	4271 $\pm$ 39	3863 $\pm$ 11	4056 $\pm$ 9	3620 $\pm$ 7	3658 $\pm$ 9	4343 $\pm$ 16	4648 $\pm$ 9
30	1	2534 $\pm$ 6	3051 $\pm$ 62	2776 $\pm$ 45	3827 $\pm$ 45	3818 $\pm$ 39	4329 $\pm$ 20	3863 $\pm$ 11	4056 $\pm$ 9	3620 $\pm$ 7	3658 $\pm$ 9	4343 $\pm$ 16	4648 $\pm$ 9
	105	1	11.0 $\pm$ 0.0	3011 $\pm$ 88	3120 $\pm$ 134	4254 $\pm$ 114	4277 $\pm$ 107	3051 $\pm$ 89	3297 $\pm$ 6	3400 $\pm$ 6	3152 $\pm$ 6	3184 $\pm$ 5	3117 $\pm$ 10
10	1	666 $\pm$ 3	2066 $\pm$ 64	1693 $\pm$ 38	2927 $\pm$ 74	2816 $\pm$ 55	3027 $\pm$ 43	3297 $\pm$ 6	3400 $\pm$ 6	3152 $\pm$ 6	3184 $\pm$ 5	3117 $\pm$ 10	3557 $\pm$ 4
	20	1242 $\pm$ 3	2067 $\pm$ 62	1571 $\pm$ 67	2840 $\pm$ 41	2835 $\pm$ 35	3123 $\pm$ 18	3297 $\pm$ 6	3400 $\pm$ 6	3152 $\pm$ 6	3184 $\pm$ 5	3117 $\pm$ 10	3557 $\pm$ 4
30	1	1745 $\pm$ 4	2247 $\pm$ 58	1722 $\pm$ 58	2826 $\pm$ 75	2839 $\pm$ 56	3283 $\pm$ 19	3297 $\pm$ 6	3400 $\pm$ 6	3152 $\pm$ 6	3184 $\pm$ 5	3117 $\pm$ 10	3557 $\pm$ 4

Table 10: **Satimage** per-cell results. $T=5,000$, $K=10$, 10 seeds.

d	r	Oracle	SPSC	SPSC-Adp	LowOFUL	VOFUL	LR-Rew	SW-Lin	D-Lin	Rst-Lin	LinUCB	LinTS	SW-LinTS
5	1	35.0 $\pm$ 2.0	1480 $\pm$ 88	1513 $\pm$ 123	2739 $\pm$ 17	2704 $\pm$ 18	2660 $\pm$ 24	655 $\pm$ 7	778 $\pm$ 6	629 $\pm$ 6	680 $\pm$ 6	2397 $\pm$ 14	1675 $\pm$ 16
55	1	18.0 $\pm$ 1.0	2439 $\pm$ 81	2015 $\pm$ 75	2913 $\pm$ 11	2911 $\pm$ 8	3450 $\pm$ 52	3051 $\pm$ 6	3201 $\pm$ 9	2779 $\pm$ 3	2848 $\pm$ 9	3403 $\pm$ 11	3781 $\pm$ 10
	10	597 $\pm$ 3	1343 $\pm$ 41	1073 $\pm$ 49	2294 $\pm$ 99	2288 $\pm$ 96	2689 $\pm$ 62	3051 $\pm$ 6	3201 $\pm$ 9	2779 $\pm$ 3	2848 $\pm$ 9	3403 $\pm$ 11	3781 $\pm$ 10
	20	1218 $\pm$ 4	1506 $\pm$ 34	1201 $\pm$ 33	2455 $\pm$ 106	2479 $\pm$ 65	2879 $\pm$ 55	3051 $\pm$ 6	3201 $\pm$ 9	2779 $\pm$ 3	2848 $\pm$ 9	3403 $\pm$ 11	3781 $\pm$ 10
	30	1725 $\pm$ 8	1900 $\pm$ 38	1635 $\pm$ 40	2641 $\pm$ 54	2772 $\pm$ 47	3085 $\pm$ 53	3051 $\pm$ 6	3201 $\pm$ 9	2779 $\pm$ 3	2848 $\pm$ 9	3403 $\pm$ 11	3781 $\pm$ 10
105	1	16.0 $\pm$ 1.0	2484 $\pm$ 59	2333 $\pm$ 108	3638 $\pm$ 7	3655 $\pm$ 12	3911 $\pm$ 47	3572 $\pm$ 8	3671 $\pm$ 9	3441 $\pm$ 6	3484 $\pm$ 10	3600 $\pm$ 9	3891 $\pm$ 7
	10	538 $\pm$ 3	1377 $\pm$ 44	1112 $\pm$ 49	2431 $\pm$ 87	2361 $\pm$ 78	2703 $\pm$ 40	3572 $\pm$ 8	3671 $\pm$ 9	3441 $\pm$ 6	3484 $\pm$ 10	3600 $\pm$ 9	3891 $\pm$ 7
	20	1098 $\pm$ 5	1454 $\pm$ 33	1182 $\pm$ 34	2229 $\pm$ 57	2346 $\pm$ 46	2697 $\pm$ 39	3572 $\pm$ 8	3671 $\pm$ 9	3441 $\pm$ 6	3484 $\pm$ 10	3600 $\pm$ 9	3891 $\pm$ 7
	30	1569 $\pm$ 5	1769 $\pm$ 32	1506 $\pm$ 33	2525 $\pm$ 42	2587 $\pm$ 80	2722 $\pm$ 25	3572 $\pm$ 8	3671 $\pm$ 9	3441 $\pm$ 6	3484 $\pm$ 10	3600 $\pm$ 9	3891 $\pm$ 7

Table 11: **Fashion-MNIST** per-cell results. $T=5,000$, $K=10$, 10 seeds.

d	r	Oracle	SPSC	SPSC-Adp	LowOFUL	VOFUL	LR-Rew	SW-Lin	D-Lin	Rst-Lin	LinUCB	LinTS	SW-LinTS
55	5	539 $\pm$ 6	3773 $\pm$ 91	3792 $\pm$ 58	4487 $\pm$ 103	4475 $\pm$ 118	4719 $\pm$ 93	4169 $\pm$ 9	4324 $\pm$ 11	3801 $\pm$ 6	3868 $\pm$ 12	5183 $\pm$ 11	5357 $\pm$ 15
	10	945 $\pm$ 5	3406 $\pm$ 73	3377 $\pm$ 92	4425 $\pm$ 126	4471 $\pm$ 142	4740 $\pm$ 59	4169 $\pm$ 9	4324 $\pm$ 11	3801 $\pm$ 6	3868 $\pm$ 12	5183 $\pm$ 11	5357 $\pm$ 15
	20	1676 $\pm$ 6	3450 $\pm$ 62	3162 $\pm$ 81	4677 $\pm$ 119	4708 $\pm$ 75	4924 $\pm$ 56	4169 $\pm$ 9	4324 $\pm$ 11	3801 $\pm$ 6	3868 $\pm$ 12	5183 $\pm$ 11	5357 $\pm$ 15
105	5	399 $\pm$ 3	3200 $\pm$ 88	3208 $\pm$ 60	3365 $\pm$ 71	3522 $\pm$ 92	3688 $\pm$ 81	4152 $\pm$ 8	4236 $\pm$ 7	4016 $\pm$ 8	4034 $\pm$ 10	4376 $\pm$ 9	4496 $\pm$ 11
	10	757 $\pm$ 6	2927 $\pm$ 76	2623 $\pm$ 84	3331 $\pm$ 42	3478 $\pm$ 51	3645 $\pm$ 65	4152 $\pm$ 8	4236 $\pm$ 7	4016 $\pm$ 8	4034 $\pm$ 10	4376 $\pm$ 9	4496 $\pm$ 11
	20	1439 $\pm$ 4	2855 $\pm$ 75	2528 $\pm$ 59	3559 $\pm$ 80	3564 $\pm$ 45	3797 $\pm$ 38	4152 $\pm$ 8	4236 $\pm$ 7	4016 $\pm$ 8	4034 $\pm$ 10	4376 $\pm$ 9	4496 $\pm$ 11
200	5	393 $\pm$ 3	2949 $\pm$ 52	2877 $\pm$ 89	3134 $\pm$ 72	2869 $\pm$ 58	2948 $\pm$ 47	4099 $\pm$ 7	4158 $\pm$ 8	4059 $\pm$ 5	4059 $\pm$ 11	4101 $\pm$ 11	4206 $\pm$ 8
	10	780 $\pm$ 4	2703 $\pm$ 52	2732 $\pm$ 101	2929 $\pm$ 43	2731 $\pm$ 63	2967 $\pm$ 26	4099 $\pm$ 7	4158 $\pm$ 8	4059 $\pm$ 5	4059 $\pm$ 11	4101 $\pm$ 11	4206 $\pm$ 8
	20	1290 $\pm$ 5	2654 $\pm$ 52	2582 $\pm$ 67	3215 $\pm$ 56	3196 $\pm$ 61	3341 $\pm$ 44	4099 $\pm$ 7	4158 $\pm$ 8	4059 $\pm$ 5	4059 $\pm$ 11	4101 $\pm$ 11	4206 $\pm$ 8

Table 12: **MNIST** per-cell results. $T=5,000$, $K=10$, 10 seeds.

d	r	Oracle	SPSC	SPSC-Adp	LowOFUL	VOFUL	LR-Rew	SW-Lin	D-Lin	Rst-Lin	LinUCB	LinTS	SW-LinTS
55	5	666 $\pm$ 5	5076 $\pm$ 88	5223 $\pm$ 109	6437 $\pm$ 152	6646 $\pm$ 62	6650 $\pm$ 116	5889 $\pm$ 16	6154 $\pm$ 12	5426 $\pm$ 10	5533 $\pm$ 9	6958 $\pm$ 13	7362 $\pm$ 20
	10	1149 $\pm$ 9	4648 $\pm$ 77	4968 $\pm$ 93	6392 $\pm$ 114	6191 $\pm$ 118	6623 $\pm$ 77	5889 $\pm$ 16	6154 $\pm$ 12	5426 $\pm$ 10	5533 $\pm$ 9	6958 $\pm$ 13	7362 $\pm$ 20
	20	2196 $\pm$ 11	4745 $\pm$ 76	4549 $\pm$ 82	6420 $\pm$ 63	6352 $\pm$ 120	6745 $\pm$ 58	5889 $\pm$ 16	6154 $\pm$ 12	5426 $\pm$ 10	5533 $\pm$ 9	6958 $\pm$ 13	7362 $\pm$ 20
105	5	582 $\pm$ 5	4658 $\pm$ 104	4844 $\pm$ 68	5184 $\pm$ 74	5305 $\pm$ 56	5807 $\pm$ 93	5992 $\pm$ 11	6073 $\pm$ 8	5749 $\pm$ 11	5811 $\pm$ 10	5892 $\pm$ 18	6222 $\pm$ 11
	10	1029 $\pm$ 5	4298 $\pm$ 96	4229 $\pm$ 157	4867 $\pm$ 96	4917 $\pm$ 46	5555 $\pm$ 58	5992 $\pm$ 11	6073 $\pm$ 8	5749 $\pm$ 11	5811 $\pm$ 10	5892 $\pm$ 18	6222 $\pm$ 11
	20	1869 $\pm$ 11	4273 $\pm$ 101	4076 $\pm$ 114	5020 $\pm$ 46	4918 $\pm$ 78	5651 $\pm$ 46	5992 $\pm$ 11	6073 $\pm$ 8	5749 $\pm$ 11	5811 $\pm$ 10	5892 $\pm$ 18	6222 $\pm$ 11
200	5	539 $\pm$ 2	4377 $\pm$ 84	4260 $\pm$ 81	4672 $\pm$ 62	4718 $\pm$ 113	4797 $\pm$ 52	5751 $\pm$ 10	5786 $\pm$ 6	5659 $\pm$ 11	5661 $\pm$ 8	5568 $\pm$ 10	5774 $\pm$ 7
	10	907 $\pm$ 5	4114 $\pm$ 66	4023 $\pm$ 113	4184 $\pm$ 63	4375 $\pm$ 75	4818 $\pm$ 42	5751 $\pm$ 10	5786 $\pm$ 6	5659 $\pm$ 11	5661 $\pm$ 8	5568 $\pm$ 10	5774 $\pm$ 7
	20	1675 $\pm$ 6	4052 $\pm$ 59	3828 $\pm$ 69	4379 $\pm$ 34	4413 $\pm$ 50	5022 $\pm$ 27	5751 $\pm$ 10	5786 $\pm$ 6	5659 $\pm$ 11	5661 $\pm$ 8	5568 $\pm$ 10	5774 $\pm$ 7

Table 13: **MovieLens** (fully real ratings) per-cell results. $T=5,000$, $K=10$, 10 seeds.

d	r	Oracle	SPSC	SPSC-Adp	LowOFUL	VOFUL	LR-Rew	SW-Lin	D-Lin	Rst-Lin	LinUCB	LinTS	SW-LinTS
55	5	4013 $\pm$ 33	6213 $\pm$ 63	6060 $\pm$ 77	7676 $\pm$ 164	7504 $\pm$ 143	8103 $\pm$ 60	6313 $\pm$ 6	6363 $\pm$ 6	6239 $\pm$ 10	6258 $\pm$ 8	6396 $\pm$ 11	6587 $\pm$ 11
	10	4802 $\pm$ 24	6250 $\pm$ 46	6015 $\pm$ 37	7689 $\pm$ 77	7461 $\pm$ 113	7855 $\pm$ 47	6313 $\pm$ 6	6363 $\pm$ 6	6239 $\pm$ 10	6258 $\pm$ 8	6396 $\pm$ 11	6587 $\pm$ 11
	20	5402 $\pm$ 21	6266 $\pm$ 42	6216 $\pm$ 44	7407 $\pm$ 56	7329 $\pm$ 108	7391 $\pm$ 30	6313 $\pm$ 6	6363 $\pm$ 6	6239 $\pm$ 10	6258 $\pm$ 8	6396 $\pm$ 11	6587 $\pm$ 11
105	5	3970 $\pm$ 33	6535 $\pm$ 56	6350 $\pm$ 73	8276 $\pm$ 126	8292 $\pm$ 89	8340 $\pm$ 35	6809 $\pm$ 5	6850 $\pm$ 4	6783 $\pm$ 4	6793 $\pm$ 6	6743 $\pm$ 13	6874 $\pm$ 12
	10	4951 $\pm$ 37	6560 $\pm$ 43	6297 $\pm$ 55	7801 $\pm$ 173	7927 $\pm$ 154	8200 $\pm$ 61	6809 $\pm$ 5	6850 $\pm$ 4	6783 $\pm$ 4	6793 $\pm$ 6	6743 $\pm$ 13	6874 $\pm$ 12
	20	5697 $\pm$ 23	6636 $\pm$ 32	6566 $\pm$ 40	7562 $\pm$ 103	7688 $\pm$ 90	7674 $\pm$ 27	6809 $\pm$ 5	6850 $\pm$ 4	6783 $\pm$ 4	6793 $\pm$ 6	6743 $\pm$ 13	6874 $\pm$ 12
200	5	3939 $\pm$ 27	6787 $\pm$ 52	6516 $\pm$ 101	8365 $\pm$ 102	8450 $\pm$ 118	8420 $\pm$ 46	7396 $\pm$ 6	7421 $\pm$ 6	7404 $\pm$ 4	7417 $\pm$ 7	7134 $\pm$ 12	7203 $\pm$ 7
	10	4998 $\pm$ 48	6844 $\pm$ 57	6576 $\pm$ 36	8383 $\pm$ 131	8163 $\pm$ 135	8513 $\pm$ 50	7396 $\pm$ 6	7421 $\pm$ 6	7404 $\pm$ 4	7417 $\pm$ 7	7134 $\pm$ 12	7203 $\pm$ 7
	20	5871 $\pm$ 25	6995 $\pm$ 47	6844 $\pm$ 53	7296 $\pm$ 118	7605 $\pm$ 117	7627 $\pm$ 58	7396 $\pm$ 6	7421 $\pm$ 6	7404 $\pm$ 4	7417 $\pm$ 7	7134 $\pm$ 12	7203 $\pm$ 7

Table 14: **Warfarin** ($d=93$, $K=8$, $T=5,000$) per-rank results across $r \in \{1, 2, 3, 5, 10\}$. 10 seeds.

d	r	Oracle	SPSC	SPSC-Adp	LowOFUL	VOFUL	LR-Rew	SW-Lin	D-Lin	Rst-Lin	LinUCB	LinTS	SW-LinTS
93	1	6.4 $\pm$ 0.7	1482 $\pm$ 49	626 $\pm$ 44	1491 $\pm$ 26	1497 $\pm$ 23	1533 $\pm$ 20	2096 $\pm$ 18	2175 $\pm$ 22	1977 $\pm$ 15	2006 $\pm$ 17	1949 $\pm$ 20	2060 $\pm$ 22
	2	93.9 $\pm$ 1.3	1404 $\pm$ 35	706 $\pm$ 37	1076 $\pm$ 130	918 $\pm$ 127	1140 $\pm$ 25	2096 $\pm$ 18	2175 $\pm$ 22	1977 $\pm$ 15	2006 $\pm$ 17	1949 $\pm$ 20	2060 $\pm$ 22
	3	177 $\pm$ 4	1372 $\pm$ 33	657 $\pm$ 39	683 $\pm$ 97	779 $\pm$ 93	1056 $\pm$ 40	2096 $\pm$ 18	2175 $\pm$ 22	1977 $\pm$ 15	2006 $\pm$ 17	1949 $\pm$ 20	2060 $\pm$ 22
	5	272 $\pm$ 7	1369 $\pm$ 30	769 $\pm$ 37	905 $\pm$ 110	930 $\pm$ 76	1178 $\pm$ 32	2096 $\pm$ 18	2175 $\pm$ 22	1977 $\pm$ 15	2006 $\pm$ 17	1949 $\pm$ 20	2060 $\pm$ 22
	10	569 $\pm$ 4	1481 $\pm$ 19	1014 $\pm$ 35	1441 $\pm$ 63	1680 $\pm$ 59	2078 $\pm$ 19	2096 $\pm$ 18	2175 $\pm$ 22	1977 $\pm$ 15	2006 $\pm$ 17	1949 $\pm$ 20	2060 $\pm$ 22

Table 15: **Open Bandit (ZOZOTOWN)** per-cell results. $T=5,000$, $K=10$, 10 seeds. Mean $\pm$ SE costed regret. Green = SPSC variant beats every non-oracle baseline; underline = best non-oracle. LinTS / SW-LinTS not run on this benchmark.

d	r	Oracle	SPSC	SPSC-Adp	LowOFUL	VOFUL	LR-Rew	SW-Lin	D-Lin	Rst-Lin	LinUCB
55	5	34.2 $\pm$ 0.5	83.0 $\pm$ 2.3	91.1 $\pm$ 4.0	126.2 $\pm$ 1.8	126.7 $\pm$ 1.5	126.8 $\pm$ 1.2	103.0 $\pm$ 0.5	106.6 $\pm$ 0.5	101.9 $\pm$ 0.6	102.2 $\pm$ 0.5
	10	65.2 $\pm$ 0.8	93.9 $\pm$ 1.5	94.1 $\pm$ 1.7	129.3 $\pm$ 2.6	130.2 $\pm$ 1.9	130.9 $\pm$ 1.3	103.0 $\pm$ 0.5	106.6 $\pm$ 0.5	101.9 $\pm$ 0.6	102.2 $\pm$ 0.5
105	5	29.7 $\pm$ 0.5	106.6 $\pm$ 4.8	103.3 $\pm$ 2.9	152.6 $\pm$ 2.3	150.1 $\pm$ 1.9	150.9 $\pm$ 1.6	107.4 $\pm$ 1.1	108.4 $\pm$ 0.4	106.4 $\pm$ 0.5	107.7 $\pm$ 1.0
	10	59.0 $\pm$ 0.7	106.9 $\pm$ 3.0	106.3 $\pm$ 3.2	153.3 $\pm$ 1.0	155.4 $\pm$ 1.8	156.1 $\pm$ 1.5	107.4 $\pm$ 1.1	108.4 $\pm$ 0.4	106.4 $\pm$ 0.5	107.7 $\pm$ 1.0




Figure 2: **Pendigits operating-regime grid.** SPSC and SPSC-Adaptive dominate every non-oracle baseline once $d \geq 55$ and $r \geq 10$.




Figure 3: **Satimage operating-regime sunburst.** Regret ratios across the (d, r) grid. Blue cells (ratio < 1) indicate SPSC wins; red cells indicate losses. (a) SPSC vs. LinUCB; (b) SPSC vs. best non-oracle competitor.

856 **Reading the tables.** Three qualitative patterns are consistent across datasets. (i) *Rank-*
857 *insensitive competitors.* LinUCB, SW-LinUCB, D-LinUCB and Restart-LinUCB operate in
858 ambient space and therefore produce the *same* regret in every row of a d block. When a
859 row shows all four of these at identical values, this is by design: it reflects the r -insensitivity
860 of ambient-rate baselines. (ii) *Subspace-aware competitors* (LowOFUL, VOFUL, LowRank-
861 Reward) do improve with d increasing but they assume a fixed subspace and therefore
862 still trail SPSC on segments where the subspace changes. (iii) *Within the synthetic and*
863 *UCI grids, SPSC-Adaptive typically trails SPSC-Alg. 1 by a small margin;* this gap reflects
864 the detector’s false-alarm budget. The ordering reverses on the clinical benchmarks (most
865 strikingly Vancomycin at $r=1$, where SPSC-Adaptive cuts regret by 37% vs. SPSC-Alg. 1;
866 see Table 19), where the oracle schedule of Alg. 1 is misaligned with the true cohort change
867 points and the detector’s adaptivity wins.

868 **Synthetic (d, r) grid.** Table 16 gives the full synthetic grid underlying Figure 1. SPSC
869 wins on 37/40 cells; losses are confined to $d \leq 10$ or $r \leq 1$ (probe-cost dominant).

Table 16: **Synthetic grid (full 40 cells).** Mean costed regret across 10 seeds for 10 methods over $(d, r) \in \{5, 10, 20, 30, 45, 60, 80, 100\} \times \{1, 3, 5, 10, 15, 20\}$ with $r < d$, $T=5,000$, $K=10$, probe period 50, window $W=400$, $\lambda=0.01$. **Methods:** SPSC (Algorithm 1, ours); **Adpt** = SPSC-Adaptive (Algorithm 2, ours); LinUCB; **Oracle** = Oracle-LinUCB (given the true subspace); **SW-Lin** = SW-LinUCB; **D-Lin** = D-LinUCB; **Rst** = Restart-LinUCB; **LR-Rew** = LowRank-Reward; LowOFUL; VOFUL. S/L = SPSC/LinUCB; S/Best = SPSC/ $\min\{\text{non-oracle, non-Adpt baselines}\}$. **Verdict.** “SPSC” if $S/L < 1$ (SPSC has lower regret than LinUCB); “LinUCB” otherwise. Bold = SPSC verdict. Counts: **37 SPSC**, **3 LinUCB**; the 3 LinUCB cells are all in the small- d corner $d \leq 10$ with r close to d , where the probe cost dominates the dimensionality savings.

d	r	SPSC	Adpt	LinUCB	Oracle	SW-Lin	D-Lin	Rst	LR-Rew	LowOFUL	VOFUL	S/L	S/Best	verdict
5	1	355	344	203	8	213	234	206	416	514	553	1.742	1.742	LinUCB
5	3	339	384	221	134	221	254	217	629	930	889	1.529	1.563	LinUCB
5	10	270	289	255	24	260	271	258	254	317	315	1.055	1.060	LinUCB
10	3	389	422	399	130	391	428	386	456	629	595	0.976	1.008	SPSC
10	5	465	525	487	254	482	523	471	660	800	823	0.955	0.988	SPSC
20	1	231	228	251	21	251	253	248	179	238	240	0.923	1.290	SPSC
20	3	353	377	456	134	449	467	451	344	449	475	0.774	1.026	SPSC
20	5	420	453	565	232	555	581	558	468	575	578	0.744	0.898	SPSC
20	10	631	699	800	498	780	840	784	804	946	975	0.789	0.809	SPSC
20	15	805	901	899	736	893	954	883	1011	1235	1222	0.895	0.912	SPSC
30	1	189	206	220	24	220	222	219	154	198	203	0.859	1.225	SPSC
30	3	307	330	407	125	409	413	404	296	355	355	0.754	1.038	SPSC
30	5	372	421	544	217	541	553	536	411	498	492	0.683	0.904	SPSC
30	10	572	615	775	451	763	790	760	690	855	828	0.738	0.829	SPSC
30	15	741	810	930	663	917	957	913	956	1066	1113	0.796	0.811	SPSC
30	20	890	960	1018	838	1002	1056	1009	1109	1236	1259	0.874	0.888	SPSC
45	1	165	169	188	25	187	189	186	130	172	168	0.881	1.275	SPSC
45	3	271	296	371	118	371	375	369	273	312	294	0.730	0.992	SPSC
45	5	360	389	504	212	506	510	502	376	419	431	0.715	0.957	SPSC
45	10	499	520	675	395	665	684	668	559	646	650	0.739	0.893	SPSC
45	15	651	684	819	577	812	836	811	736	811	808	0.796	0.885	SPSC
45	20	811	836	973	758	954	980	952	901	988	989	0.833	0.900	SPSC
60	1	141	146	165	30	166	166	165	108	134	134	0.858	1.311	SPSC
60	3	229	243	296	117	294	296	292	222	249	239	0.776	1.033	SPSC
60	5	315	319	414	199	412	416	414	320	366	343	0.760	0.985	SPSC
60	10	441	479	605	367	600	608	601	536	564	585	0.729	0.822	SPSC
60	15	588	605	743	524	740	752	738	713	752	736	0.791	0.824	SPSC
60	20	708	753	853	663	845	868	845	838	882	862	0.830	0.845	SPSC
80	1	110	125	118	27	118	118	118	91	111	106	0.931	1.214	SPSC
80	3	231	250	290	116	292	293	290	224	238	239	0.796	1.034	SPSC
80	5	279	295	381	187	379	384	379	293	307	319	0.733	0.954	SPSC
80	10	419	430	542	345	539	545	538	447	493	473	0.772	0.937	SPSC
80	15	534	576	671	486	668	671	667	587	638	655	0.796	0.911	SPSC
80	20	640	663	767	599	761	770	761	691	744	761	0.835	0.927	SPSC
100	1	107	111	112	28	113	113	113	79	86	93	0.954	1.359	SPSC
100	3	192	194	238	107	236	238	236	176	193	201	0.809	1.094	SPSC
100	5	272	273	346	183	343	346	344	266	294	290	0.787	1.024	SPSC
100	10	368	388	471	317	472	472	470	403	438	440	0.780	0.912	SPSC
100	15	466	491	573	430	571	576	572	517	530	535	0.813	0.902	SPSC
100	20	609	634	726	572	725	733	723	661	692	702	0.839	0.921	SPSC

870 G Sensitivity and robustness studies

871 G.1 Probe-rate ablation

872 Sweeping the probe period over $\{5, 10, 20, 30, 50, 100, 300\}$ at $d=4$, $r=1$, $K=4$, $T=6,000$
 873 produces a clearly U-shaped regret curve (Figure 4): sparse probing under-recovers the
 874 subspace; over-frequent probing sacrifices exploitation. Optimum at probe period 20–30,
 875 matching $m_k^* \propto \ell_k^{2/3}$ prediction. The right panel shows a monotone relationship between
 876 late-segment subspace error and final regret, directly linking SPSC’s gains to subspace
 877 recovery quality.




Figure 4: Probe-rate ablation. Left: final regret vs. probe frequency. Right: final control regret vs. late-segment subspace error.

878 G.2 Rank misspecification

879 On Covertypes ($d=155$, true effective rank $r^*=10$, $K=4$, $T=10,000$, 5 seeds), sweeping spec-
 880 ified r at the grid points $\{1, 3, 5, 10, 15, 20, 30, 50, 80\}$ (min $r=1$, max $r=80$) gives the U-
 881 shaped curve in Figure 5. Underestimation loses signal (+33% at $r=1$); mild overestima-
 882 tion captures residual structure (sweet spot at $r=30$, −10%). SPSC beats LinUCB across
 883 $r \in [15, 50]$, so mild overestimation is a safe default.




Figure 5: Rank misspecification robustness (Covertypes, $d=155$, $r^*=10$). SPSC beats LinUCB for $r \in [15, 50]$; sweet spot at $r=30$ (10% improvement).

884 G.3 Robustness and necessity experiments (A/B/C)

885 Three ablations on the reference setting ($d=4$, $r=1$, $K=4$, $T=6,000$, probe period 30,
886 $W=100$, 10 seeds).

887 **A. Variance misspecification.** Injecting $|\delta_\sigma| \in \{0, 0.01, 0.02, 0.05, 0.1, 0.2, 0.5\}$ into the
888 probe statistic $s_t = y_t^2 - (\sigma_\varepsilon^2 + \delta_\sigma)$. Regret and subspace error are stable through $|\delta_\sigma| = 0.5$
889 (more than $5\times$ the true variance $\sigma_\varepsilon^2 = 0.09$); at maximum misspecification, regret changes
890 by $< 5\%$. For scaled-sphere probes (the theory probe distribution), a constant centering
891 error contributes the population-level identity shift $\tilde{B} = -(\delta_\sigma/d)I_d$, which uniformly shifts
892 eigenvalues without corrupting eigenvectors.

893 **B. Bounded cross-correlation.** Injecting $\varepsilon_t \rightarrow \varepsilon_t + \epsilon_\times x_t^\top \theta_t$ for $\epsilon_\times \in$
894 $\{0, 0.01, 0.05, 0.1, 0.2, 0.5, 1.0\}$. Even at $\epsilon_\times = 1.0$ (cross-correlation equal to signal), re-
895 gret increases by $< 1\%$, indicating that moderate violations of orthogonality are practically
896 benign.

897 **C. Imperfect probe coverage.** Restricting probe directions to the first $d_{\text{cov}} \in \{1, 2, 3, 4\}$
898 coordinates produces a dramatic monotonic regret blow-up: at $d_{\text{cov}} = 1$, SPSC regret rises
899 to 4,289, nearly matching LinUCB (4,463); subspace error degrades to 0.75. This directly
900 confirms the necessity results: restricted coverage destroys identifiability and forces $\Omega(T)$
901 regret within the unobserved directions.




Figure 6: Robustness and necessity experiments. (a) Variance misspecification, regret stable through $|\delta_\sigma| = 0.5$. (b) Cross-correlation relaxation, $< 1\%$ degradation through $\epsilon_\times = 1.0$. (c) Imperfect coverage, regret blow-up with restricted probe directions.

902 G.4 Assumption-violation master table (large-scale)

903 We complement the small-scale A/B/C ablations above with a large-scale ($d=60$, $r=5$,
904 $K=10$, $T=5,000$) assumption sweep over (a) variance misspecification $\hat{\sigma}^2/\sigma_\varepsilon^2 \in \{0.5, 1, 2, 4\}$,
905 (b) approximate-rank perturbation $\epsilon_k^\perp \in \{0, 0.05, 0.1, 0.2, 0.4\}$, and (c) spectral radius
906 $\rho(A_k) \in \{0.8, 0.95, 0.99, 1.0\}$ (the last violates the stable-LDS hypothesis). Variance mis-
907 spec is absorbed into the scaled-identity bias (Lemma B.6) and the relative gap stays flat
908 ($\hat{\sigma}^2/\sigma_\varepsilon^2$ ratio $1.27 \rightarrow 1.16$ as the sweep goes $0.5 \rightarrow 4$); approximate-rank degrades regret
909 near-linearly in $\epsilon_k^\perp$; SPSC remains well-behaved through $\rho = 0.99$ and degrades gracefully
910 at $\rho = 1$ (random-walk limit, outside the stable-LDS assumption).

911 G.5 Subspace recovery rate and change-point adaptation

912 Figure 7 plots the binned mean subspace error $\|\hat{P}_k - P_k^*\|_2$ vs. probe count; the empirical rate
913 matches the predicted $1/\sqrt{m}$ on a log-log scale. Figure 8 plots post-change instantaneous
914 regret and the subspace-error trajectory: SPSC recovers substantially faster than ambient
915 LinUCB after each boundary.

916 G.6 Noise robustness, change-point frequency, drift speed

917 Three single-parameter sweeps on the reference setting.

Table 17: **Assumption-violation master table** ($d=60, r=5, K=10, T=5,000$, 10 seeds, costed dynamic regret, mean $\pm$ SE). SPSC tracks Oracle through all three sweeps; bold = best non-Oracle in row.

Sweep	Level	Oracle	SPSC-Alg1	SPSC-Adap	LinUCB	SW-LinUCB	D-LinUCB	LowOFUL
$\hat{\sigma}^2/\sigma_\epsilon^2$	0.5	221 $\pm$ 5	281 $\pm$ 10	304 $\pm$ 11	397 $\pm$ 17	395 $\pm$ 16	401 $\pm$ 18	370 $\pm$ 19
	1.0	257 $\pm$ 6	315 $\pm$ 10	319 $\pm$ 10	413 $\pm$ 18	412 $\pm$ 18	415 $\pm$ 19	358 $\pm$ 17
	2.0	299 $\pm$ 10	357 $\pm$ 13	355 $\pm$ 13	421 $\pm$ 19	421 $\pm$ 18	423 $\pm$ 20	377 $\pm$ 14
	4.0	344 $\pm$ 12	398 $\pm$ 16	398 $\pm$ 16	426 $\pm$ 19	425 $\pm$ 19	427 $\pm$ 20	390 $\pm$ 20
$\epsilon_k^\perp$	0.00	257 $\pm$ 6	315 $\pm$ 10	319 $\pm$ 10	413 $\pm$ 18	412 $\pm$ 18	415 $\pm$ 19	358 $\pm$ 17
	0.05	255 $\pm$ 7	311 $\pm$ 12	331 $\pm$ 10	414 $\pm$ 18	414 $\pm$ 18	417 $\pm$ 19	379 $\pm$ 20
	0.10	261 $\pm$ 7	319 $\pm$ 13	331 $\pm$ 12	422 $\pm$ 19	422 $\pm$ 18	425 $\pm$ 19	376 $\pm$ 20
	0.20	286 $\pm$ 8	340 $\pm$ 13	348 $\pm$ 6	452 $\pm$ 19	453 $\pm$ 18	457 $\pm$ 20	399 $\pm$ 19
	0.40	358 $\pm$ 9	396 $\pm$ 10	410 $\pm$ 10	559 $\pm$ 21	561 $\pm$ 20	565 $\pm$ 22	497 $\pm$ 13
$\rho(A_k)$	0.80	713 $\pm$ 6	745 $\pm$ 7	747 $\pm$ 9	744 $\pm$ 8	743 $\pm$ 6	745 $\pm$ 8	732 $\pm$ 9
	0.95	1083 $\pm$ 11	1291 $\pm$ 16	1307 $\pm$ 14	1403 $\pm$ 16	1406 $\pm$ 15	1404 $\pm$ 19	1330 $\pm$ 18
	0.99	1132 $\pm$ 23	2017 $\pm$ 46	1929 $\pm$ 34	2751 $\pm$ 46	2732 $\pm$ 44	2750 $\pm$ 52	2585 $\pm$ 54
	1.00	914 $\pm$ 19	2899 $\pm$ 83	2813 $\pm$ 59	4644 $\pm$ 90	4523 $\pm$ 87	4674 $\pm$ 90	5552 $\pm$ 156




Figure 7: **Subspace recovery.** $\|\hat{P}_k - P_k^*\|_2$ vs. probe count matches the predicted $1/\sqrt{m}$ rate on a log-log scale.




Figure 8: **Change-point adaptation.** Instantaneous regret drops sharply after each boundary as fresh probes rebuild the subspace.

918 **Noise.** Sweeping $\sigma_\varepsilon \in \{0.05, 0.10, 0.20, 0.30, 0.50, 0.80\}$: SPSC/LinUCB ratio grows mildly
 919 with noise (from 0.566 to 0.595); SPSC/Oracle stable at $\approx 1.55\times$. The low-rank benefit is
 920 not a low-noise artifact.

921 **Change-point frequency.** $K \in \{1, 2, 4, 6, 8, 12\}$ at $T=6,000$. Ratio grows from 0.414
 922 (long segments) to 0.855 (very short); SPSC retains an edge at all frequencies but the margin
 923 shrinks as the per-segment probe budget collapses below the identifiability threshold.

924 **Drift speed.** LDS spectral radius $\rho \in \{0.30, \dots, 0.99\}$. Sharp phase transition:
 925 SPSC/LinUCB ≈ 1.01 for $\rho \leq 0.70$, drops to 0.861 at $\rho=0.95$, and to 0.582 at $\rho=0.99$.
 926 Subspace error is nearly constant across ρ (0.28–0.33), so the missing benefit at low ρ is a
 927 signal-energy effect, not estimation failure: low-rank structure in a signal already dominated
 928 by noise leaves little to extract.




Figure 9: Noise robustness: SPSC advantage is noise-monotone.




Figure 10: Change-point frequency.

929 H Comparison with nonstationary baselines (Russac benchmark)

930 **Russac et al. NeurIPS 2019 benchmark.** $d=2, r=1, K=4, T=6,000, \sigma_\varepsilon=1$, 50 arms,
 931 30 seeds. SPSC reduces control regret by 46.5% vs. D-LinUCB and 34.8% vs. SW-LinUCB;
 932 probe overhead is a 0.9% fraction.

933 I Warfarin: random-subspace ablation

934 To decompose SPSC's Warfarin gain into (i) dimensionality reduction and (ii) the *learned*
 935 subspace itself, we replace $\hat{U}_t$ with a fresh random rank- r projection every segment and
 936 run SPSC otherwise unchanged. Results on $d=93$ across $r \in \{1, 2, 3, 5, 10\}$ (10 seeds): the




Figure 11: Drift speed. Benefit gated by LDS correlation time $\tau_\rho = 1/(1 - \rho)$.

Table 18: Russac et al. benchmark. Final control regret and costed regret (mean $\pm$ SE, 30 seeds).

Algorithm	Control	Costed	vs. D-LinUCB
SPSC Alg. 1 (ours)	$2,229 \pm 54$	2,249	-46.5%
Oracle LinUCB	$1,757 \pm 55$	1,757	-57.8%
D-LinUCB [Russac et al., 2019]	$4,167 \pm 110$	4,167	—
SW-LinUCB [Cheung et al., 2019]	$3,417 \pm 82$	3,417	-18.0%
OFUL [Abbasi-Yadkori et al., 2011]	$4,645 \pm 149$	4,645	+11.5%
Reset-LinUCB (oracle)	$4,647 \pm 147$	4,647	+11.5%




Figure 12: Russac et al. NeurIPS 2019 benchmark.

random-subspace variant retains about half of the SPSC-vs-LinUCB gap, confirming that *dimensionality reduction explains roughly half* the win and the learned subspace explains the rest. The learned subspace captures 2–6 $\times$ more signal variance than a random rank- r projection at every r , directly quantifying the value of the probe-based identification.




Figure 13: **Warfarin random-subspace ablation.** SPSC’s learned subspace vs. a fresh random rank- r projection at every segment. Half of the gap to LinUCB is dimensionality reduction; the other half is learned signal.

J BOSS/Jedra adaptation to the piecewise setting

BOSS [Duong et al., 2024] and Jedra [Jedra et al., 2024] are stationary low-rank methods. To give them a fair comparison to SPSC we restart each at the *true* segment boundaries (an oracle advantage we do not give SPSC) and otherwise leave their hyperparameters at the authors’ recommended defaults. The benchmark is piecewise LDS, $K=10$, $T=5,000$, $\sigma_\varepsilon=0.3$, 40 actions, 10 seeds, probe cost $c=0.1$, on a grid $d \in \{55, 105, 200\} \times r \in \{5, 10, 20\}$ (eight cells: (55, 5), (55, 10), (105, 5), (105, 10), (105, 20), (200, 5), (200, 10), (200, 20)).

Results (cf. Table 4): SPSC Alg. 1 wins every cell, by 5.8–19.9%, with the lead largest at small d (where the ambient d is closest to r and probing pays off most) and narrowing to 5.8–8.0% at $d=200$ as oracle-restart-plus-PCA closes the gap. SPSC-Adaptive (no oracle restarts) also beats BOSS and Jedra on every cell, and the ranking is consistent on every seed. The point is that “low-rank exploitation” alone is not enough for changing-subspace problems: the learner must *identify a changing subspace online*, which is SPSC’s core capability.

K Vancomycin: full per-rank tables

This appendix gives the per-rank Vancomycin results ($d=93$, $K=8$, $T=5,000$, $\sigma_\varepsilon=0.3$, 10 seeds, probe period 10, window 200, probe cost $c=0.1$) summarized in Table 3 of the main body. SPSC-Adaptive beats every non-oracle baseline at every rank.

L When the adaptive variant wins

This appendix disentangles the two SPSC variants: Alg. 1 assumes oracle change points; Alg. 2 detects them online via the detector statistic (15). When the oracle boundaries are correct, Alg. 1 is slightly better because the adaptive detector pays a small statistical price for unknown change-point localization. When the oracle provides misspecified (false-alarm) boundaries, the picture reverses sharply.

SPSC-Adaptive’s regret is flat as K_{oracle} grows, while Alg. 1 degrades monotonically and is 3 $\times$ worse at $K_{\text{oracle}}=50$: a detector-based reset is strictly safer than a fixed schedule when boundaries are unknown or noisy.

Table 19: **Vancomycin per-rank results** (mean $\pm$ SE, 10 seeds). “vs. best comp.” compares SPSC-Adaptive against the strongest non-oracle baseline in each column (marked *). SPSC-Adaptive wins every rank.

Method	$r=1$	$r=2$	$r=3$	$r=5$	$r=10$
Oracle LinUCB	34.3 ± 1.5	115.1 ± 2.0	197.2 ± 3.0	268.1 ± 3.4	373.2 ± 5.7
SPSC Alg. 1 (ours)	840.3 ± 24.5	805.9 ± 15.4	802.3 ± 12.3	805.3 ± 9.5	811.0 ± 12.2
SPSC-Adaptive	525.3 ± 45.8	613.6 ± 33.1	640.3 ± 33.0	665.9 ± 27.2	705.6 ± 19.5
LowOFUL	1049.9 ± 12.8	$740.1 \pm 36.3^*$	734.5 ± 37.8	791.5 ± 26.0	801.9 ± 20.1
VOFUL	1063.5 ± 5.5	789.7 ± 24.8	$712.5 \pm 25.5^*$	$763.0 \pm 30.6^*$	$758.6 \pm 27.9^*$
LowRank-Reward	1175.6 ± 11.6	806.8 ± 11.2	835.1 ± 10.7	898.8 ± 6.7	880.5 ± 5.5
SW-LinUCB	875.6 ± 6.7	875.6 ± 6.7	875.6 ± 6.7	875.6 ± 6.7	875.6 ± 6.7
D-LinUCB	925.3 ± 5.8	925.3 ± 5.8	925.3 ± 5.8	925.3 ± 5.8	925.3 ± 5.8
Restart-LinUCB	$808.9 \pm 5.3^*$	808.9 ± 5.3	808.9 ± 5.3	808.9 ± 5.3	808.9 ± 5.3
LinUCB	825.3 ± 6.1	825.3 ± 6.1	825.3 ± 6.1	825.3 ± 6.1	825.3 ± 6.1
SPSC-Adaptive vs. LinUCB	-36.4%	-25.7%	-22.4%	-19.3%	-14.5%
SPSC-Adaptive vs. best comp.	-35.1%	-17.1%	-10.1%	-12.7%	-7.0%




Figure 14: **SPSC-Alg. 1 vs. SPSC-Adaptive**. Top: d -sweep with correct oracle boundaries. Bottom: oracle quality. With $K_{\text{real}}=5$ true boundaries, Alg. 1 is fed $K_{\text{oracle}} \in \{5, \dots, 50\}$ (extras are false alarms) and degrades; Alg. 2 ignores the oracle and is flat.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction list four contributions (identification, algorithm, regret bound, experiments) which are stated formally in Theorem 2.2, Algorithms 1–2, Theorem 4.1, and §5; the scope and limitations (e.g. approximate-low-rank settings, oracle vs. detected change points) are flagged in the relevant sections.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Limitations are stated where they arise: the regret bound is conditional on the windowed projected potential condition (§4); the bound is most informative when within-segment variation V_{in} is small (Theorem 4.1 discussion); the rule-of-thumb crossover $d - r \gtrsim T^{1/6}$ is qualitative, not sharp (§4, §5.1); SPSC-Adaptive is evaluated empirically only with no formal regret guarantee (§4).

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: All theoretical results (Theorem 2.2, Theorem 4.1, Corollary 4.3, Corollary 4.4, Proposition 4.5) state their assumptions explicitly; full proofs are in Appendix B together with the supporting lemmas (matrix Bernstein, projector concentration, projected nonstationary UCB lemma, prefix concentration).

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: All datasets, dimensions d , ranks r , segment counts K , horizon T , seed counts, probe rate, regularisation, window length, and per-baseline hyperparameters are listed in §5 and Appendix E; per-cell tables in Appendix F report mean $\pm$ SE for every dataset, (d, r) cell, and method.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [\[Yes\]](#)

Justification: The supplemental archive contains all experiment scripts (one per figure and per table), per-dataset environment code, cached JSON results, and a README mapping each figure and table to its generator script.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [\[Yes\]](#)

Justification: §5 states the headline parameters per benchmark; Appendix E gives the full hyperparameter table (probe period, ridge λ , window W , confidence δ , probe cost c , baseline-specific knobs) used uniformly across all (d, r) cells without per-dataset retuning.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Every reported number is $\text{mean} \pm \text{SE}$ or $\text{mean} \pm \text{SEM}$ over at least 10 (typically 10–30) seeds; figures show ± 1 SE bands; the seed count and aggregation choice are stated in each table caption and in §5.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix E reports that the full grid (all benchmarks, 10 seeds, 10 methods) takes 6–20 hours on a single 16-core CPU; no GPU is required (NumPy/SciPy only).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work is a theoretical and empirical study of a stochastic bandit algorithm on synthetic and publicly available benchmarks; the clinical environments are semi-synthetic surrogates calibrated to published cohort statistics, with no patient-identifiable data, no human subjects, and no clinical deployment.

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: The paper studies a foundational sequential-decision algorithm with no specific deployment; the positive impact is improved sample efficiency for low-rank decision problems (recommendation, dosing) and negative-impact pathways are not direct. Clinical bandit benchmarks are evaluation surrogates, not deployable dosing policies.

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases an algorithm and experiment scripts, not pre-trained models, generative models, or scraped datasets; no high-risk-of-misuse assets are distributed.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

1067 Justification: All datasets used (UCI Covertypes, Pendigits, Satimage; MNIST;
1068 Fashion-MNIST; MovieLens-100K; Open Bandit Dataset [Saito et al., 2021]) and
1069 all baseline algorithms (LinUCB, D-LinUCB, SW-LinUCB, LowOFUL, VOFUL,
1070 LowRank-Reward, LinTS, BOSS, Jedra) are cited at first use; clinical benchmarks
1071 follow the IWPC [Consortium, 2009] and Rybak et al. [Rybak et al., 2020] dosing
1072 guidelines.

1073 **13. New assets**

1074 Question: Are new assets introduced in the paper well documented and is the
1075 documentation provided alongside the assets?

1076 Answer: [Yes]

1077 Justification: The supplemental material releases the SPSC and SPSC-Adaptive
1078 Python implementations, the per-benchmark environments, all experiment scripts,
1079 and a README documenting the figure-to-script and table-to-source mappings.

1080 **14. Crowdsourcing and research with human subjects**

1081 Question: For crowdsourcing experiments and research with human subjects, does
1082 the paper include the full text of instructions given to participants and screenshots,
1083 if applicable, as well as details about compensation (if any)?

1084 Answer: [N/A]

1085 Justification: The paper does not involve crowdsourcing or research with human
1086 subjects.

1087 **15. Institutional review board (IRB) approvals or equivalent for research
1088 with human subjects**

1089 Question: Does the paper describe potential risks incurred by study participants,
1090 whether such risks were disclosed to the subjects, and whether Institutional Review
1091 Board (IRB) approvals (or an equivalent approval/review based on the requirements
1092 of your country or institution) were obtained?

1093 Answer: [N/A]

1094 Justification: The paper does not involve human subjects; no IRB approval is re-
1095 quired.

1096 **16. Declaration of LLM usage**

1097 Question: Does the paper describe the usage of LLMs if it is an important, original,
1098 or non-standard component of the core methods in this research? Note that if the
1099 LLM is used only for writing, editing, or formatting purposes and does *not* impact
1100 the core methodology, scientific rigor, or originality of the research, declaration is
1101 not required.

1102 Answer: [N/A]

1103 Justification: LLMs are not part of the core methodology; the contributions (al-
1104 gorithm, identification result, regret bound, experiments) are independent of any
1105 LLM.